
Formal Languages and Compilers

Jordan Pyott

2022-01-21

Contents

Course Information	3
Information about these notes	3
Course Staff	3
Assessments and Grading	3
Textbooks/Resources	3
Lectures	4
Introduction	4
Finite Automata and Regular Languages	5
Symbols, Strings and Languages	5
Deterministic Finite Automata	8
Non-Deterministic Finite Automata	11
Regular Expressions	16
Non-Regular Languages	20
Modelling Independent Processes	23
Context-Free Languages	23
Context-Free Grammars	23
Pushdown automata (Stack automaton)	27
Compilers	28
How does a compiler convert these instructions?	29
Lexical Analysis	29
Scanners and Automata	30
Syntax Trees	31

Course Information

Information about these notes

These notes are meant to be used in conjunction to lecture material provided from COSC261 at the University of Canterbury, there are references and links to particular lectures and videos about processes that will not be available to those reading this later

Course Staff

- Coordinator/Lecturer
 - Walter Guttman
 - * 033692451
 - * walter.guttman@canterbury.ac.nz

Assessments and Grading

Grading policy

1. You must achieve an average grade of at least 50% over all assessment items
2. You must achieve an average grade of at least 45% over all invigilated assessments

Assessment Items

- Quizzes (15%)
- Assignment Superquiz (10%)
- Lab Test (35%)
- Final Exam (40%)

Textbooks/Resources

No textbooks are required, but see the following book for additional information:

- Carol Critchlow and David Eck; Fundamentals of Computation; version 2.3.1, 2011

Lectures

Introduction

Topic overview of the course

- pattern matching
 - Regular expressions describe patterns
 - Search using REGEX is supported in many programs
 - Can all patterns be described by regular expressions?
 - One to one with state diagrams and automata
- Compilers
 - Programs can be run by an interpreter or by compiling them first
 - Interpreting may be slow
 - Compiling to machine code avoids much of the overhead
 - Compiler performs analysis, code generation and optimisation
 - How can these tasks be automated for different programming languages?
- Syntax analysis
 - Analyses code to determine if the syntax is correct for compiling, this is done by using context free grammars *there are other methods of doing this however this is the main one we will look at in this course*
 - *does the syntax conform to the languages grammar?*
 - Ideally we want to generate the parser for our language, we will look into how to manually like are parser and how regular expressions and pattern matching can be used to evaluate this behaviour.
- Code generation
 - There are formalisms that exist to generate code in order to create compilers via code generation.

Finite Automata and Regular Languages

Symbols, Strings and Languages

Languages

- An alphabet Σ is non-empty finite set of *symbols*/
- A string over Σ is a finite sequence of symbols from Σ
- The length of $|\sigma|$ of a string σ is the number of symbols in σ
- The empty string ϵ is the unique set of length 0
- Σ^* is the set of all strings over Σ
- A language L over Σ is a set of strings $L \subseteq \Sigma^*$

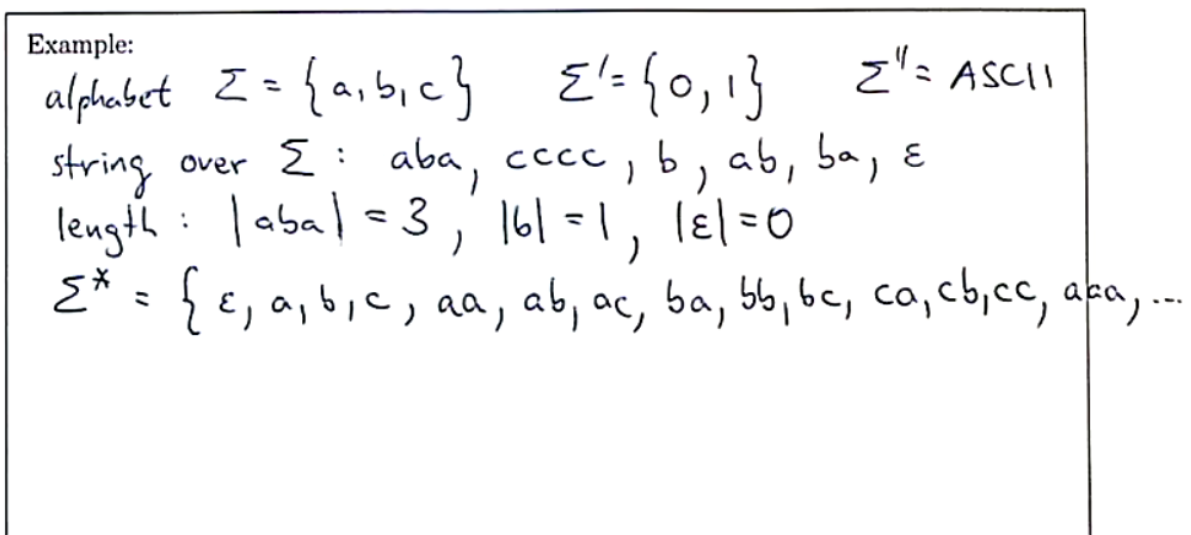


Figure 1: Example

- Note that with a finite alphabet we can have an infinite size for Σ^*
 - This is because we have not specified a size for our length of elements within Σ^*

An example of a set that we might use is the unicode set as Sigma.

For example:

- $\text{python} \subset \text{UNICODE}$
- $\text{english} \subset \text{UNICODE}$

Because of this relationship, we can use filtering, searching and [REGEX](#) in order to manipulate and set rules around this relationship (or syntax in the case of programming languages by using comparisons and combination of formalisms).

Let $a, b \in \Sigma$ be symbols and let $x, y, z \in \Sigma$ be strings. - Symbols and strings can be concatenated by writing one after the other - xy is the concatenated version of x and y . - Note that concatenation is accociative - ϵ is an identity for concatenation $\epsilon x = x = x\epsilon$ - $|xy| = |x| + |y|$

Lifting to a set

Let $A, B \subseteq \Sigma$ be languages:

- concatenate languages A and B by conciliating each string from A with each from B
- $AB = \{xy | x \in A, y \in B\}$
- Language concatenation is associative
- $\{\epsilon\}$ is the identity of language concatenation

$$- x^0 = \epsilon$$

Example: $L_1 = \{\epsilon, ab, abb\}$ $L_2 = \{a, ba\}$

$L_1 L_2$	a	ba
ϵ	a	ba
ab	aba	$abba$
abb	$abba$	$abbbba$

$L_2 L_1$	ϵ	ab	aba
a	a	aab	$aaba$
ba	ba	$baab$	$baaba$

 $L_1 L_2 = \{a, ba, aba, abba, abbbba\}$

Figure 2: Lifting Example

Concatenation can be iterated

- a^n is the string comprising n copies of the symbol $a \in \Sigma$
- x^n is the string that concatenates n copies of the string $x \in \Sigma$

- These operations are defined *inductively*
- The base case is $x^0 = \epsilon$
- The *inductive case* is $x^{n+1} = x^n x$

Example: $a^5 = aaaaa$

A^n is defined similarly for a language $A \subseteq \Sigma^*$.

$$* A^0 = \{\epsilon\}.$$

$$* A^{n+1} = AA^n.$$

$$* A^1 = A, A^2 = AA, A^3 = AAA, \dots$$

$$* A^* = \bigcup_{n \in \mathbb{N}} A^n = A^0 \cup A^1 \cup A^2 \cup A^3 \cup \dots$$

$$* A^* = \{x_1 x_2 \dots x_{n-1} x_n \mid x_i \in A \text{ for each } 1 \leq i \leq n \text{ for some } n \in \mathbb{N}\}.$$

$$* A^+ = \bigcup_{n \geq 1} A^n = AA^* = A^1 \cup A^2 \cup A^3 \cup \dots$$

Figure 3: Powers of Language

Take aways, the * symbol means that we have zero or more of something, + means that we have one or more of something (this is how we use this notation practically in regex expressions)

Key notation and definitions

- Sets are languages
- variables are strings
- variables with index are symbols

Deterministic Finite Automata

A *deterministic finite automaton (DFA)* is a structure $M = (Q, \Sigma, \delta, q_0, F)$ where:

- Q is a non-empty finite set, the *states*,
- Σ is a non-empty finite set, the *input alphabet*
- $\delta : Q \times \Sigma \rightarrow Q$ is the *transition function*
- $q_0 \in Q$ is the *start state*,
- $F \subseteq Q$ is the set of *accept states* or *final states*.

This can be shown in a *transition diagram* for a visual indication of how this might work:

see lecture 3, 14:00 minutes for information on how to construct these transition diagrams

Deterministic Finite Automata Example Ex 1: The following DFA accepts strings over $\{a, b\}$.

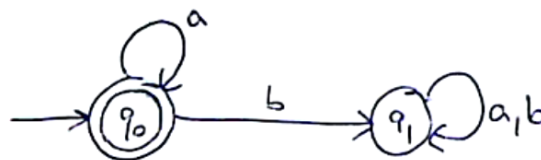


Figure 4: Sink state diagram

Note, we are trying to contain an automata that is as least complicated as possible (least nodes). In the above example, we have this concept of a true state and a false state to the specified condition. We can also note that state q_1 is a *sink state* as once we reach q_1 it is impossible to leave that state (*this is because the condition specifies that we need to contain 0 b's*).

Ex 2: DFA accepting all strings over $\{a, b\}$ with a number of a -symbols that is not a multiple of 4.

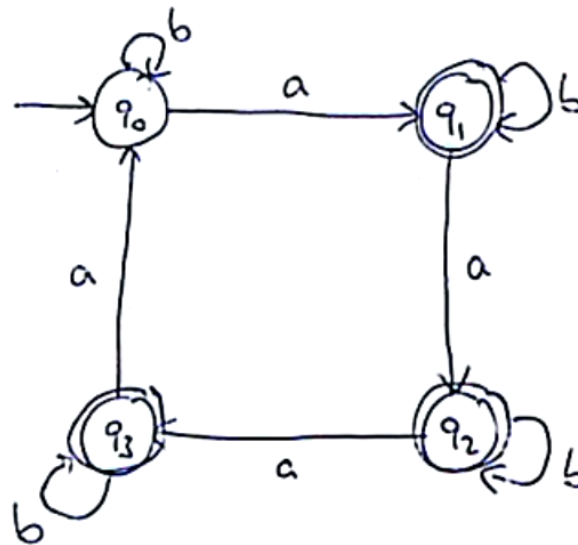


Figure 5: Modulo four example

Closure of properties of Deterministic Finite Automata Exptended transition function

$$\hat{\delta} : Q \times \Sigma^* \rightarrow Q$$

$$\hat{\delta}(q, \epsilon) = q$$

$$\hat{\delta}(q, ax) = \hat{\delta}(\delta(q, a), x) \quad \text{where } a \in \Sigma, x \in \Sigma^*$$

Now we have seen the basics of DFA's, so how do we get from these basic models to something that is more complicated? We can use a combination of DFA's in order to define more complicated systems.

Regular languages are closed under:

- complement
 - All strings that are not in the set, (generic approach to create such an automaton)
 - We can get the complement by swapping accepting and non-accepting states
 - If we have a regular language, we know that the complement will always be regular
 - * Formal proof provided in [Lecture Four](#): 27:10
- intersection or product automaton
 - We can combine two automata using an *intersection*, we can use set theory in order to satisfy two automata at the same time. We want to check with one automaton to see if the string is satisfied

- $X \cap Y$ is regular if both X and Y are regular
 - Accept states are defined as an acceptance of **both** automata, not just one
 - The product automaton accepting the intersection of the two languages is (*synchronous*):
 - * Example found in [Lecture Four](#): 41:00
- union
 - concatenation
 - star

Non-Deterministic Finite Automata

The following automaton accepts strings with a symbol 1 in the third position from the end. It is not a DFA because there are two 1-transitions in state q_0 and no transitions in state q_3 .

- **DFA defined by:** $\delta : Q \times \Sigma \rightarrow Q$
- **NFA defined by:** $\delta : Q \times \Sigma \rightarrow P(Q)$
 - where $P(Q) = \{S \mid S \subseteq Q\}$, $P(Q)$ is also called the *power set*

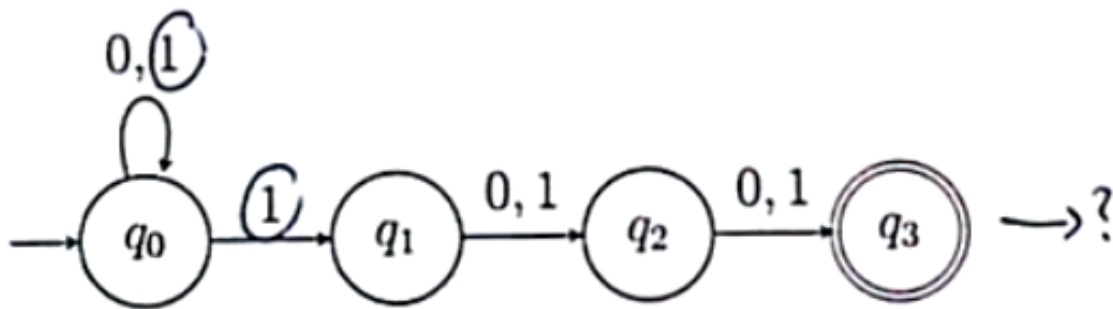


Figure 6: Non-Deterministic Finite Automata

Here is an example of the above transition relation:

δ	0	1
q_0	$\{q_0\}$	$\{q_0, q_1\}$
q_1	$\{q_2\}$	$\{q_2\}$
q_2	$\{q_3\}$	$\{q_3\}$
q_3	$\{\emptyset\}$	$\{\emptyset\}$

Possible transition sequences for input 1101 are:




Figure 7: Possible transition states of this non-deterministic finite automata

The extended transition relation $\hat{\delta} : Q \times \Sigma^* \rightarrow P(Q)$ is

- $\hat{\delta}(q, \epsilon) = \{q\}$
- $\hat{\delta}(q, ax) = \bigcup_{p \in \delta(q,a)} \hat{\delta}(p,x)$ where $a \in \Sigma$ and $x \in \Sigma^*$

Example of evaluating this extension can be found in [Friday March 4th Lecture: 8:00](#)

The result of this will tell you what states are available from the current state (in a recursive nature)
 If we check the acceptance criteria for the automaton, we can do this by using this method to get to the last state.

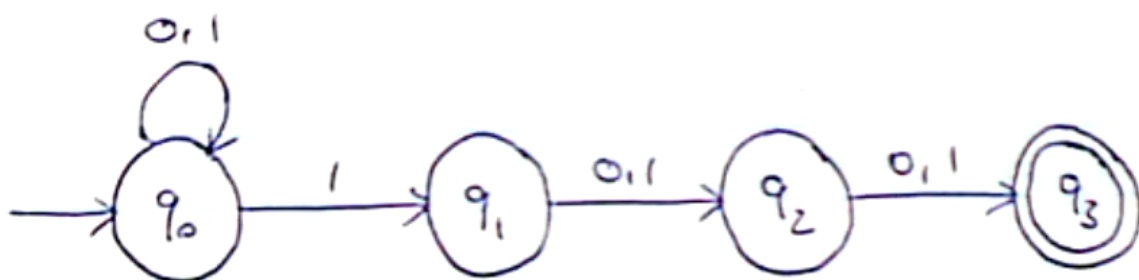


Figure 8: Consider this Automaton

The Subset Construction We need to consider all the possible nodes that we can reach from the current state. Below is an example of how we might do such a traverse on the above automaton. Note that we don't know what state we are going to be in, we are considering all possible options at a point in time.

In the following example, the numbers in each node are shorthand notation for the set containing all possible nodes that we can reach from that value, hence $q_{012} \rightarrow \{q_0, q_1, q_2\}$.

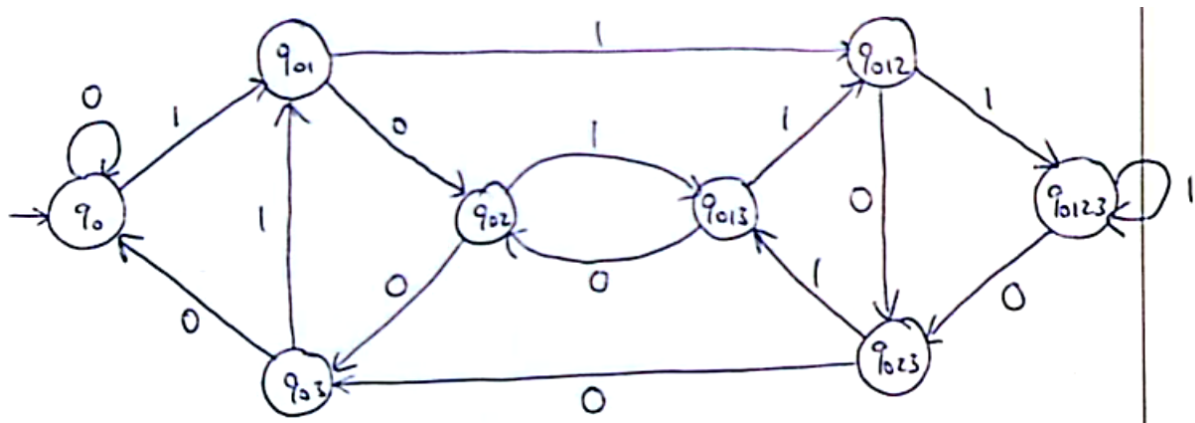


Figure 9: Subset Construction

The above construction considers all possibilities of the **non-deterministic automaton**, this allows us to construct a **deterministic finite automata**, the only change we need to make is to add acceptance states to this construction.

Every language accepted by an **NFA** is accepted by a **DFA**

Proof: NFA $M = (Q, \Sigma, \delta, q_0, F)$

idea: keep track of all states M can be in while reading the input

subset automaton (DFA) $M' = (\mathcal{P}(Q), \Sigma, \delta', \{q_0\}, F')$

$$F' = \{ S \subseteq Q \mid S \cap F \neq \emptyset \}$$

$$\delta'(S, a) = \bigcup_{q \in S} \delta(q, a)$$

$$\Rightarrow \hat{\delta}'(S, w) = \bigcup_{q \in S} \hat{\delta}(q, w) \quad \otimes$$

To show $L(M') = L(M)$. For every $w \in \Sigma^*$,

$$w \in L(M')$$

$$\stackrel{(\Rightarrow)}{\text{Def. DFA } L(M')} \hat{\delta}'(\{q_0\}, w) \in F'$$

$$\stackrel{(\Rightarrow)}{\text{Def. } F'} \hat{\delta}'(\{q_0\}, w) \cap F \neq \emptyset$$

$$\stackrel{(\Rightarrow)}{\otimes} \left(\bigcup_{q \in \{q_0\}} \hat{\delta}(q, w) \right) \cap F \neq \emptyset$$

$$\stackrel{(\Rightarrow)}{\text{Def. } \cup} \hat{\delta}(q_0, w) \cap F \neq \emptyset$$

$$\stackrel{(\Rightarrow)}{\text{Def. NFA } L(M)} w \in L(M)$$

Through this lecture, we have found an equivalence relation between DFA's and NFA's, meaning that we can convert DFA's to NFA's, therefore NFA's accept exactly regular languages. The number of states may grow exponentially in the subset construction.

Non-Deterministic Finite Automata with ϵ -Transitions The following automaton accepts the union of two regular languages. It is not a DFA because of the ϵ -transitions in state r_0 .




Figure 10: Epsilon transitions

An NFA with ϵ transitions is a structure $M = (Q, \Sigma, \delta, q_0, F)$ where:

- Q, Σ, q_0 and F are as in a DFA,
- ϵ is a special symbol with $\epsilon \notin \Sigma$,
- $\delta : Q \times (\Sigma \cup \{\epsilon\}) \rightarrow P(Q)$ is the transition relation
- δ may have ϵ -transitions and yields a set of successor states

This is a way to decouple and clarify a set of states and a choice of two transitions in an automaton, this allows us to clearly identify and mark unions in the automaton.

Here is an example of a NFA that is designed in order to check validity of different types of numbers for a programming language. *Hexadecimal, decimal, octal and binary.*

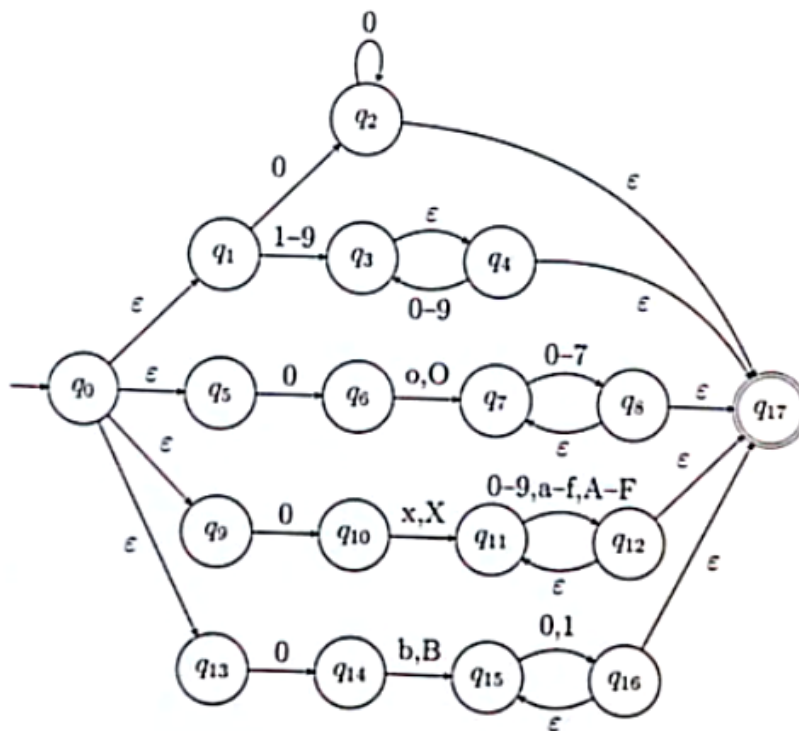


Figure 11: numerical automaton

Regular Expressions

These are patterns that can be used to match substrings in a given string:

- `ls *201?.*` lists files whose name without extension ends in 201 followed by some character
- `ls *a*a*a*` lists any file whose name contains three `a` characters
- `rm *.log` deletes all log files
- `grep '[A-Z][a-z]\{3,7\}'` finds lines with a capital followed by 3-7 lower case letters

Atomic patterns are:

- $a \forall a \in \Sigma$ is matched by the symbol a
- ϵ is matched by the empty string
- θ is matched by nothing
- $?$ is matched by any symbol in Σ

Compound patterns are formed from patterns p and q as follows:

- $p|q$ is matched by string w if w matches p or q .

- pq is matched by w if $w = xy$ and x matches p and y matches q .
- $\neg p$ matches by w if w does not match p .
- $[p]$ is matched by w if w is empty or matches p .

Here are some examples of a compound expression that can be evaluated in order to get its set of mapped values.

$$L((a|b)) = L(a) \cup L(b) = \{a\} \cup \{b\} = \{a, b\}$$

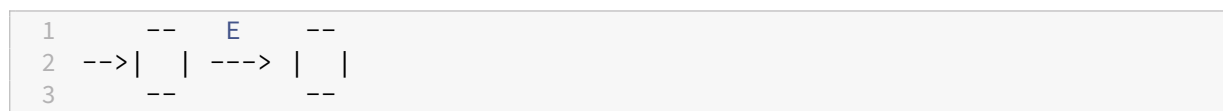
$$L((a|b)(c|a)) = L(a) \cup L(b) \cup L(c) \cup L(a) = \{a, b\} \cup \{c, a\} = \{ac, aa, bc, ba\}$$

In order to preserve readability, we can omit parentheses in order to simplify the expressions. Here is an example: $p|qr^* = (p|(q(r)^*))$

It is important to realise that we can construct **ANY** regular expression, and all such expressions will be accepted by a corresponding NFA.

Creating Automata to only accept single constructs

For ϵ construct:



For \emptyset construct:



More examples can be found in [Lecture 8](#): 30:00

NOTE: include all standard constructs for automata that are commonly used

The above automata can be used for formulation of more complicated regular expressions, we simplify these **at your own risk** as we may be able to save some states (making program more efficient) however we may find that we are leaving out cases that should be covered by our automaton.

Using Regular Expressions as Symbols within Automata

- Transitions may be labelled with regular expressions instead of just symbols from Σ or ϵ
- There is just one accept state
- There is at most one transition between any two states

- There are no transitions into the start state or out of the accept state




Figure 12: Regex Transitions Example

We will see how we can translate any automaton into a regular expression. The main idea behind how we achieve this is to try to eliminate states from the diagram.

Every language accepted by an NFA is generated by a regular expression. Proof:

- * Idea: successively eliminate states from generalised NFA.
- * Done if just the start state and the accept state are left.
- * Otherwise choose a state q_j that is neither the start state nor the accept state.
- * Make the following change for all states q_i and q_k with $i \neq j$ and $k \neq j$:

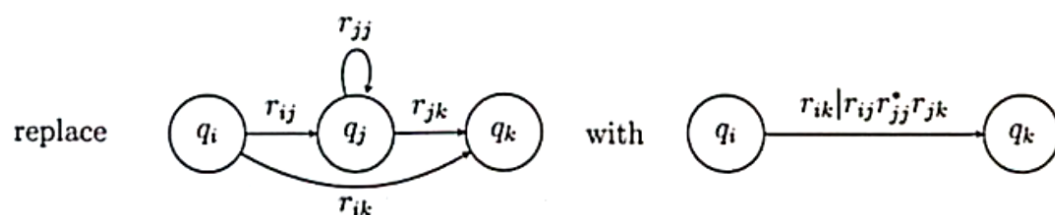


Figure 13: Converting automaton to regex

The hard part of the above process is that we must apply this algorithm to all possible nodes and edges, this means that we must apply the same pattern to all possible ways (this is unusual and some times difficult).

Minimisation of Deterministic Finite Automata

A method to minimise the number of states in a DFA is:

1. Eliminate states which cannot be reached from the start state
2. Find equivalent states
3. Collapse equivalent states

The two states are *equivalent* if:

- $\hat{\delta}(p, w) \in F \iff \hat{\delta}(q, w) \in F \forall w \in \Sigma^*$
- The automaton accepts the same strings when started in p or in q .
- Collapsing p and q does not change the accepted language
- p and q are *distinguishable* if they are not equivalent

Distinguishable states can be obtained as follows:

- Any $p \in F$ and $q \notin F$ are distinguishable by $w = \epsilon$
- Let $\delta(p, a) = r$ and $\delta(q, a) = s$ for $a \in \Sigma$
 - If r and s are distinguishable by $(w = x)$ then p and q are distinguishable by $(w = ax)$

Equivalence in finite automata

We use the notation $p \sim q$ if states p and q are equivalent, the relation \sim has:

- is reflexive: $p \sim p \quad \forall \quad p \in Q$
- is symmetric: $p \sim q$ implies $q \sim p$ for all $p, q \in Q$
- is transitive $p \sim q$ and $q \sim r$ implies $p \sim r$ for each $p, q, r \in Q$

An *equivalence relation* is a relation $\sim \subseteq A \times A$ that is reflexive, symmetric and transitive.

- $[a] = \{b \in A \mid a \sim b\}$ is the equivalence class of $a \in A$
- a is representative of its equivalence class $[a]$
- $A / \sim = \{[a] \mid a \in A\}$ is the *quotient* of A by \sim
 - This is the set of all equivalence classes

Automaton can start from any nodes that are found in the equivalence class that starting node s_0 is in.

Note that we minimal DFA's are unique up to isomorphism

The minimisation algorithm makes a quotient construction:

- * Let DFA $M = (Q, \Sigma, \delta, q_0, F)$ be given.
- * The *quotient automaton* is $M' = (\underbrace{Q/\sim}, \Sigma, \delta', [q_0], F')$.
- * A state of M' is an equivalence class of \sim .
- * $F' = \{[q] \mid q \in F\}$ comprises the equivalence classes of the accept states of M .
- * $\delta' : \underbrace{Q/\sim} \times \Sigma \rightarrow \underbrace{Q/\sim}$ is defined by $\delta'([q], a) = [\delta(q, a)]$.
- * δ' is well-defined because $[\delta(q, a)]$ does not depend on which representative q is chosen.

Figure 14: Algorithm to minimise quotient

Minimisation algorithm Quotient construction

Decision Problems for Regular Languages Simple yes or no problems about the properties of a given set

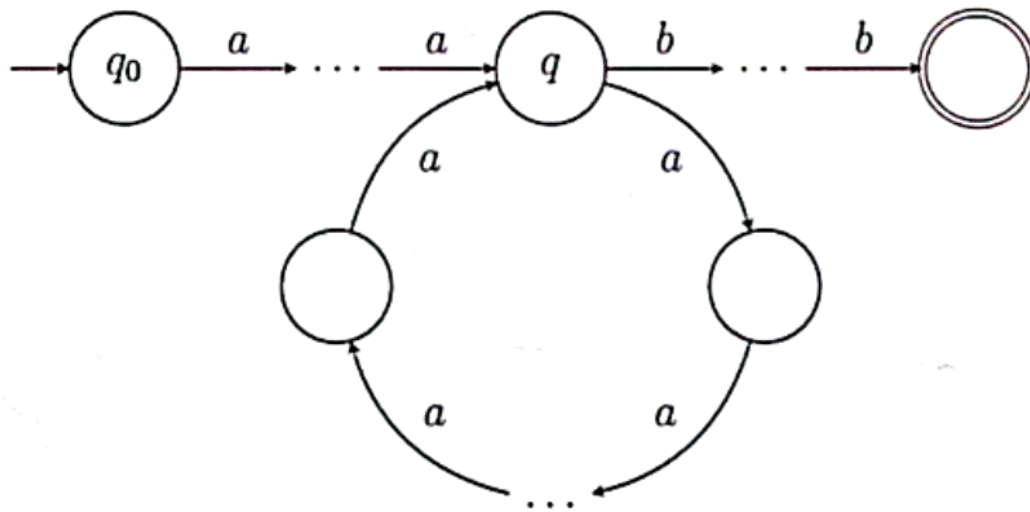
See Lecture 10, 24:00 minutes to see how these are done Note, Go through these examples and actually implement them for end of year exams []

Non-Regular Languages

The language $A = a^n b^n \mid n \in \mathbb{N}$ is a non-regular language.

The idea behind this proof is that we can construct a cycle where some iterations of the cycle are not included in the acceptance state.

- Assume that A is regular, we can prove this by applying contradiction
 - $A = L(M)$ for DFA $M = (Q, \Sigma, \delta, q_0, F)$ with k states
- The transition sequence for input a^k contains $k + 1$ states
- By the pigeonhole principle, a state q is visited twice

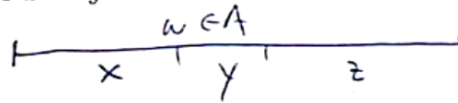


- There are i, j with $- \leq i \leq j \leq k$ and $\hat{\delta}(q_0, a^i) = \hat{\delta}(q_0, a^j)$
- Hence $\hat{\delta}(q_0, a^i b^i) = \hat{\delta}(q_0, a^j b^i)$
- But $a^i b^i \in A$ and $a^j b^i \notin A$, so $a^i b^i \in F$ and $a^j b^i \notin F$
- This is a contradiction, so the assumption does not hold

The Pumping Lemma The pumping lemma can be used to identify a contradiction to be used in a proof.

Pumping Lemma: Let A be regular. Then there is a number n such that each $w \in A$ with $|w| \geq n$ can be broken into three parts $w = xyz$ with

- * $y \neq \varepsilon$,
- * $|xy| \leq n$, and
- * $xy^iz \in A$ for each $i \in \mathbb{N}$.



$xz \in A$
 $xy^2z \in A$
 $xy^3z \in A$
 \vdots

Use the pumping lemma to show that $A = \{a^n b^n \mid n \in \mathbb{N}\}$ is not regular:

Assume A is regular.

Let $n \in \mathbb{N}$.

Set $w = a^n b^n$.

Let x, y, z be such that $w = xyz$ and $y \neq \varepsilon$ and $|xy| \leq n$. $\leadsto y = a^j$ for $j \geq 1$

Set $i = 2$.

Then $xy^iz \in A$ by the pumping lemma, but $xy^2z = xy^2z = a^{n+j}b^n \notin A$. Contradiction.

The pumping lemma is an exchange of information between a set of agents, we are able to obtain certain values from this lemma in order to obtain more information about the problem we are trying to prove.

Sorting a set

\forall input : input is a set \Rightarrow

\exists output : $\text{set}(\text{output}) = \text{input} \wedge$
 $\text{len}(\text{output}) = |\text{input}| \wedge$

$\forall i : 0 < i < \text{len}(\text{output}) \Rightarrow$
 $\text{output}[i-1] \leq \text{output}[i]$

Figure 15: Sorting a list example

The above figure is displaying how we might solve a programming problem using mathematics and the pumping lemma, the take away of this is that when we are solving problems in computer science, we are essentially proving a mathematical statement. If we can prove the method from this construction, we do not need to test the output, as we can guarantee that if the proof holds and is the same as the problem construction, then the program has to work with 100% consistency.

The pumping lemma shows us that a few languages are non-regular, and if we can map other languages to the pumping lemma's non-regular languages, then we can determine that that language is also non-regular

Modelling Independent Processes

This is a model of concurrency within a program. The *shuffle* operation describes two independent processes:

- The shuffle $x \parallel y$ if two strings $x, y \in \Sigma^*$ contains all possible interleavings of their symbols
- For example, $ab \parallel cd = \{abcd, acbd, acdb, cabd, cadb, cdab\}$
- Shuffle of strings is defined inductively:

$$\begin{aligned} \varepsilon \parallel y &= \{y\} \\ x \parallel \varepsilon &= \{x\} \\ ax \parallel by &= \{a\}(x \parallel by) \cup \{b\}(ax \parallel y) \quad \text{for } a, b \in \Sigma \end{aligned}$$

Figure 16: Shuffle of strings

Context-Free Languages

Context-Free Grammars

A *context-free grammar* (CFG) is a structure $G = (N, \Sigma, P, S)$ where

- N is a finite set, the *non-terminals*
- Σ is a finite set disjoint from N , the *terminals*
- $P \subseteq N \times (N \cup \Sigma)^*$ is a finite set of *Productions*
- $S \in N$ is the *start symbol*

Productions are denoted as follows:

- A production $(A, w) \in P$ is written $A \rightarrow w$
- Several productions are written as $A \rightarrow w_1 | \dots | w_n$
- The right-hand side may be empty: an ϵ -production is written $A \rightarrow \epsilon$

Here is an example of G that matches with arithmetic expressions:

$$\Sigma = \{+, *, (,), n\}$$

$$N = \{E, T, F\}$$

$$S = E$$

$$P = E \rightarrow T | E + T$$

$$T \rightarrow F | T * F$$

$$F \rightarrow (E) | n$$

Note that the production notation can be expressed as a piecewise

$$E = \begin{cases} E \rightarrow T \\ E \rightarrow E + T \end{cases}$$

This is the same expression as shown above in the example

Using the previous example, we will see how we can compare strings using these CFG's, see Lecture: [Tue Mar 22 : 16 Minutes](#).

Chomsky Normal Form A CFG $G = (N, \Sigma, P, S)$ is in chomsky normal form if every production has the form:

- $A \rightarrow BC$ where $B, C \in N$ or,
- $A \rightarrow a$ where $a \in \Sigma$
- The right hand side of each production is either two non-terminals or a terminal

For every CFG G with $\varepsilon \notin L(G)$ there is a CFG G' in Chomsky normal form with $L(G) = L(G')$.

1. Eliminate ε -productions of the form $A \rightarrow \varepsilon$.
 2. Eliminate unit-productions of the form $A \rightarrow B$.
 3. Eliminate non-generating non-terminals.
 4. Eliminate non-reachable non-terminals.
 5. Eliminate terminals from right-hand sides of length at least 2.
 6. Eliminate right-hand sides of length at least 3.
1. If $A \rightarrow uBv \in P$ for $u, v \in (\Sigma \cup N)^*$ and $B \rightarrow \varepsilon \in P$, add $A \rightarrow uv$ to P .
 - * Repeat this step while there are changes.
 - * Afterwards, remove all ε -productions of the form $A \rightarrow \varepsilon$.
 2. If $A \rightarrow B \in P$ and $B \rightarrow w \in P$ for $w \in (\Sigma \cup N)^*$, add $A \rightarrow w$ to P .
 - * Repeat this step while there are changes.
 - * Afterwards, remove all unit-productions of the form $A \rightarrow B$.
 3. A non-terminal A is *generating* if $A \Rightarrow^* w$ for some $w \in \Sigma^*$.
 - * If $A \rightarrow w \in P$ and w contains only terminals, A is generating.
 - * If $A \rightarrow w \in P$ and w contains only terminals or generating non-terminals, A is generating.
 - * Remove each non-generating non-terminal with all productions containing it.
 4. A non-terminal A is *reachable* if $S \Rightarrow^* uAv$ for some $u, v \in (\Sigma \cup N)^*$.
 - * S is reachable.
 - * If $A \rightarrow w \in P$ and A is reachable, every non-terminal in w is reachable.
 - * Remove each non-reachable non-terminal with all productions containing it.
 5. Consider every terminal a in a right-hand side of length at least 2.
 - * Add a new non-terminal A and the production $A \rightarrow a$.
 - * Replace every occurrence of a in a right-hand side of length at least 2 with A .
 6. Consider every production $A \rightarrow B_1B_2 \dots B_n$ with $n \geq 3$.
 - * Add a new non-terminal C and replace this production with $A \rightarrow B_1C$ and $C \rightarrow B_2B_3 \dots B_n$.
 - * Repeat this step while there are changes.

Figure 17: Steps for converting to Chomsky Normal Form

An example of how to convert to Chomsky Normal Form is run through in [Fri Mar 25 Lecture: 20 Minutes](#)

The Cocke-Younger-Kasami Algorithm This is a recursive algorithm to solve the membership problem $w \in L(G)$, we start by assuming that G is in chomsky normal form, then we let n be the length of w . Mark the indices for the symbols in w and then compute between positions i, j until we generate w_{ij} . The end result is calculating $N_{ij} = \{A \in N \mid A \rightarrow^* w_{ij}\}$.

Given a string $w \in \Sigma^*$ and a CFL A , is $w \in A$?

- * This is the test for *membership* in a CFL.
- * It is similar to parsing, but no syntax tree has to be constructed.
- * $A = L(G) = \{x \in \Sigma^* \mid S \Rightarrow_G^* x\}$ for a CFG G .
- * Hence $w \in A$ if and only if $S \Rightarrow_G^* w$.
- * Checking all derivations does not work, since there might be infinitely many.
- * Assume that G is in Chomsky normal form.
- * Every non-terminal produces at least one terminal.
- * It suffices to consider derivations that introduce up to $|w|$ non-terminals.
- * This gives an upper bound on the length of derivations that need to be checked.
- * The number of derivations might still be exponential in the length of w .

This is a technique to improve running time, it is a specific use of *dynamic programming* (bottom up approach). It works by creating smaller problems in order to solve bigger problems and avoids repeating calculation of previous smaller problems.

Decision Problems For Context-Free Languages

- Membership $x \in \Sigma^*$
 - Convert G into chomsky normal form and apply CYK
- Emptiness $L(G) = \emptyset$
 - Check if the start symbol is non-generating
- Finiteness $|L(G)| \in \mathbb{N}$
 - Convert G into Chomsky normal form
 - Construct a graph with:
 - * Nodes: non-terminals
 - * edges: for every $A \rightarrow BC$ add an edge (A, B) and an edge (A, C)
- Universality $L(G) = \Sigma^*$ and intersection emptiness $L(G) \cap L(G') = \emptyset$:

- There are no algorithms to solve this problem
- Equivalence $L(G) = L(G')$
 - There is no algorithm to solve this problem
- Inclusion $L(G) \subseteq L(G')$
 - There is no algorithm to solve this problem

Closure properties of Context-Free Languages CFL's are closed under union:

- Let $L_1 = L(G_1)$ and $L_2 = L(G_2)$
- Assume that $N_1 \cap N_2 = \emptyset$ and let $S \notin N_1 \cup N_2$.
- Construct $G = (N_1 \cup N_2 \cup \{S\}, \Sigma, P_1 \cup P_2 \cup \{S \rightarrow S_1, S \rightarrow S_2\}, S)$
- Then $L(G) = L_1 \cup L_2$

Similar constructions show that CFL's are closed under concatenation and star. The intersection of a CFL and a regular language is a context free languages. CFL's are **not** closed under intersection.

Pushdown automata (Stack automaton)

A pushdown automaton is an extension of a finite automaton

- A stack is added in which stack symbols can be stored
- Transitions may depend on symbols at the top of the stack
- Transitions may change the symbols at the top of the stack
- The stack size is unlimited

PDA definition

- A PDA is some structure: $M = Q, \Sigma, \Gamma, q_0, F$
- Γ is a finite set, the stack *alphabet*
- $\delta : Q \times (\Sigma \cup \{\epsilon\}) \times \Gamma^* \rightarrow P(Q \times \Gamma^*)$ is the transition relation

The transition relation has the following properties:

- As in NFA's with ϵ transitions, there are two sources of non-deterministic_finite_automata
- δ may have ϵ transitions
- δ yields a set of successor states
- Extending NFA's the transitions refer to the top symbol of the stack

The goal of PDA's Acceptance

This is a little bit different to NFA's and DFA's, the goal of a PDA is that the stack must be empty at the end of a string in order to be accepted. this means that anything that is added must be removed in order for the automaton to accept.

Configuration sequences:

- * \rightarrow^* is the reflexive-transitive closure of \rightarrow .
- * $c \rightarrow^* d$ means configuration d is reachable from configuration c in zero or more transitions.
- * The language accepted by M is

$$L(M) = \{x \in \Sigma^* \mid (q_0, x, \varepsilon) \rightarrow^* (q, \varepsilon, \varepsilon) \text{ for some } q \in F\}.$$

- * Acceptance requires the PDA to be in an accept state and to have an empty stack.

Figure 18: PDA Acceptance criteria

For this type of automaton, it is worthy to note that the states are not as important as the transitions, we often can reduce data to very few states with this type of automaton.

Notes about PDA's

- We cannot take an arbitrary PDA and execute an algorithm to get a deterministic PDA that accepts the same language. *This is impossible.*
- We can generate an equivalent context free language for every given PDA.

Compilers

Interpreters and compilers:

- Assume a high level program P takes input I and produces output O .
- An interpreter takes P and I and produces O .
- Compiling breaks this process into two parts.
- The compiler takes P and produces low-level machine code M .
- The processor takes M and I and produces O .
- Matched syntax or strings within a language are called **Token's**

Languages like Java have both an interpreter and a compiler. There is no reason why you cannot do this for any given language, it is merely a matter of work. Typically compilers go right down to machine code. An example of this is the C language, it does this by using two steps. One to convert C code to assembly and then a second compiling step in order to convert assembly to raw machine code.

How does a compiler convert these instructions?

We can use expression/syntax trees in order to model the parser. This allows us to see what is happening in the design under the hood.

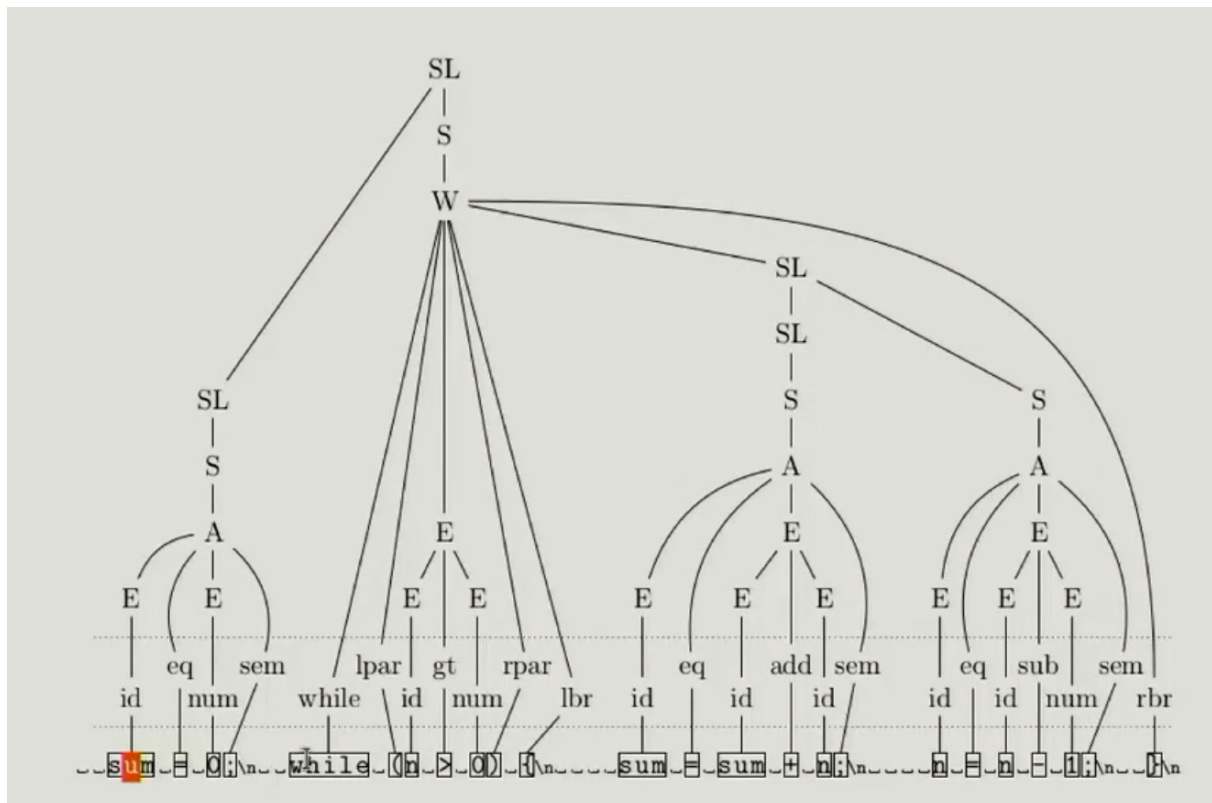


Figure 19: Here is a small overview of how this works

These trees can be optimised in order to speed up the process. Most of these optimisations can be performed before it is taken down to the target language. The general way this is translated is as follows:

- The machine holds the values of variables in a small number of *registers*
- Calculations take place on the *stack*
- The contents of registers are pushed and popped from the stack
- The processor executes the code step by step unless instructed to jump

Lexical Analysis

The first step of compiling is performed by the scanner:

- It reads the source code character by character
- It produces a sequence of tokens to be fed to the parser
- White-space and comments are discarded

There are several kinds of tokens:

- reserved words
- symbols
- identifiers (such as variable names)
- number constants
- string constants
- Identifiers and constants that have concrete values attached

Each of these are described or defined by a regular expression, for reserved words and symbols this might be just a simple string. For other kinds, a sequence of definitions might be used.

When defining these tokens, it is important to follow the following rules:

- Such a sequence of definitions must not be cyclic
- It can be converted into a flat regular expression by repeated substitution
- Character classes abbreviate choices
- The resulting regular expression is converted to a DFA that is used for matching

Use of the scanner:

- The scanner will not convert the entire source code at once.
- Typically it is called by the parser to deliver the next token in the source.
- The parser needs *lookahead* that shows the next token without consuming it.

Scanners and Automata

The scanner recognises several kinds of tokens at the same time. This allows us to group identifiers, symbols and reserve words into groups. To design this, we can construct an NFA in order to easily combine the given automata (using epsilon transitions connecting to each of the converted NFA's for each of the rules).

Once we have our big automaton, we can convert it to a DFA and then minimise in order to get an automaton that we can use.

In order to take the complement of an regex, we can construct a regex, convert it to a DFA and take the complement for a given automaton.

In the lectures [Wed Apr 6](#) there is an example on how to build a DFA for a C-style comment for a programming language. This can be found at 19:00.

Notes

Terminals used in context free grammars are defined as the tokens found in a programming language. The non-terminals associated with this context are [Expressions](#), [Statements](#), [Comparisons](#) as these are the generators of tokens.

Syntax Trees

The parser constructs a syntax tree from a sequence of tokens

- Every leaf of the syntax tree is labelled with a terminal
- Every inner node of the syntax tree corresponds with the application of a rule
- The inner node is labelled with a non-terminal on the left hand side
- The tree is ordered: the sequence of children matters

We can see how to construct these syntax trees in the [Fri Apr 8](#) lecture at time 24:00. There are two different methods to do this construction. We can use the Top down approach or the bottom up approach (both of these approaches are outlined in this section of the lectures).