
COSC367: Artificial Intelligence

This course introduces major concepts and algorithms in Artificial Intelligence. Topics include problem solving, reasoning, games, and machine learning.

Jordan Pyott

2021-05-06 16:57

Contents

Artificial Intelligence	3
Course Information	3
Textbooks / Resources	3
Readings	4
Lectures	4
Lecture One: Searching the State Space	4
Lecture Two: Searching the State Space (part two)	12

Artificial Intelligence

Course Information

The course covers core topics in AI including:

- uninformed and informed graph search algorithms,
- propositional logic and forward and backward chaining algorithms,
- declarative programming with Prolog,
- the min-max and alpha-beta pruning algorithms,
- Bayesian networks and probabilistic inference algorithms,
- classification learning algorithms,
- consistency algorithms,
- local search and heuristic algorithms such as simulated annealing, and population-based algorithms such as genetic search and swarm optimisation.

Grades

Standard Computer science policy applies

- Average 50% over all assessment items
- Average at least 45% on all invigilated assessment items

Grading structure for course

- Assignments (5%)
 - Two Super Quiz's
- Quizzes (16.5%)
 - Weekly Quiz Assessments (1.5% ea)
- Lab Test (20%)
- Final Exam (58.5%)

Textbooks / Resources

- Poole, David L. 1958, Mackworth, Alan K; Artificial intelligence : foundations of computational agents; Cambridge University Press, 2010.
- Russell, Stuart J, Norvig, Peter; Artificial intelligence : a modern approach; 3rd ed; Prentice Hall, 2010.

Readings

Lectures

Lecture One: Searching the State Space

What is state?

- A state is a data structure that represents a possible configuration of the world *agent and environment*
- The **state space** is the set of all possible states for that problem
- actions change the state of the world
- Example: A vacuum cleaner agent in two adjacent rooms which can be either clean or dirty.

- Location = {left, right}
- Left-room-condition = {dirty, clean}
- Right-room-condition = {dirty, clean}
- State-space = Location
 × Left-room-condition
 × Right-room-condition

In this example, each state is represented by a triple (3-tuple).

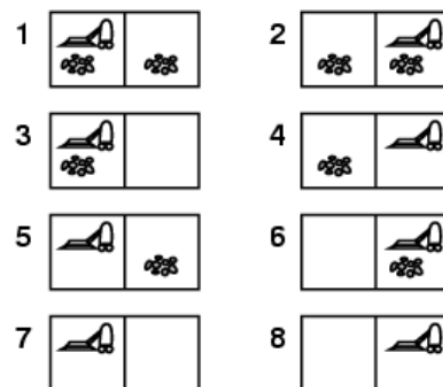


Figure 1: State space example one

State can also be represented as a graph *both directed and undirected*

- Example: Suppose the vacuum cleaner agent can take the following actions: L (go left), R (go right), S (suck).

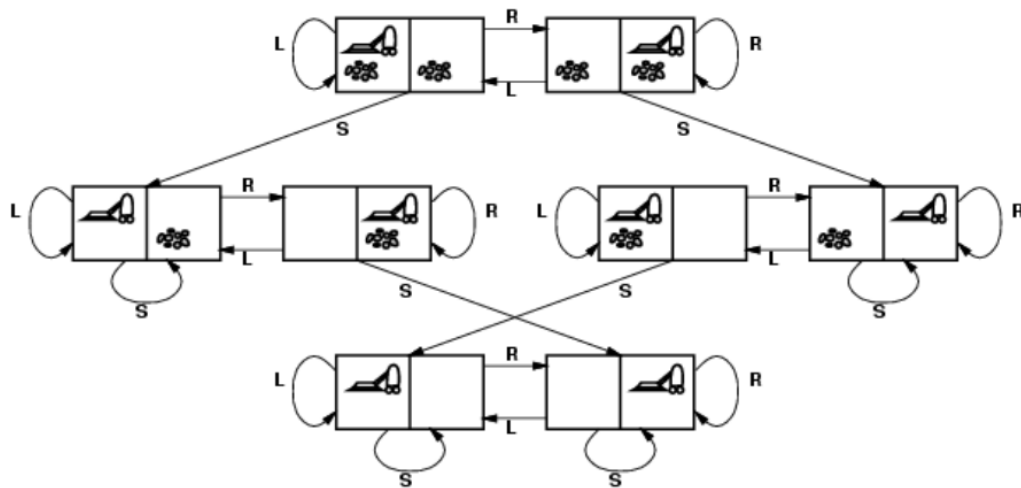


Figure 2: State space graph simplified

- Many problems in AI can be abstracted to the problem of finding a path in a directed graph
- Notation we use is **Nodes** and **arcs** for **vertices** and **edges** in a graph

Explicit vs Implicit graphs

- In **explicit graphs** nodes and arcs are readily available, they are read from the input and stored in a data structure such as an adjacency list/matrix.
 - the entire graph is in memory.
 - the complexity of algorithms are measured in the number of nodes and/or arcs.
- In **implicit graphs** a procedure `outgoing_arcs` is defined that given a node, returns a set of directed arcs that connect node to other nodes.
 - The graph is generated as needed *due to the complexity of the graphs*.
 - The complexity is measured in terms of the depth of the goal state node or *how far do we have to get into the graph to find a solution*.

Explicit graphs in quizzes

- In some exercises we use small explicit graphs to study the behaviour of various frontiers
- Nodes are specified in a set

- Edges are specified in a list
 - pairs of nodes, or triples of nodes (in a tuple)

Searching graphs

- We will use generic search algorithms: given a graph, start nodes, and goal nodes, incrementally explore paths from the start nodes.
- Maintain a **frontier** of paths that have been explored
 - frontier: paths that we have already explored
- As search proceeds, the frontier is updated and the graph is explored until a goal node is found.
- The order in which paths are removed and added to the frontier defines the search strategy
- A **search tree** is a tree drawn out of all the possible actions in terms of a tree.
 - How do we handle loops? *Covered in next lecture*
 - In the search tree outlined below, you can see that the *end of paths on frontier* represents a BFS relationship note this is not always the case.

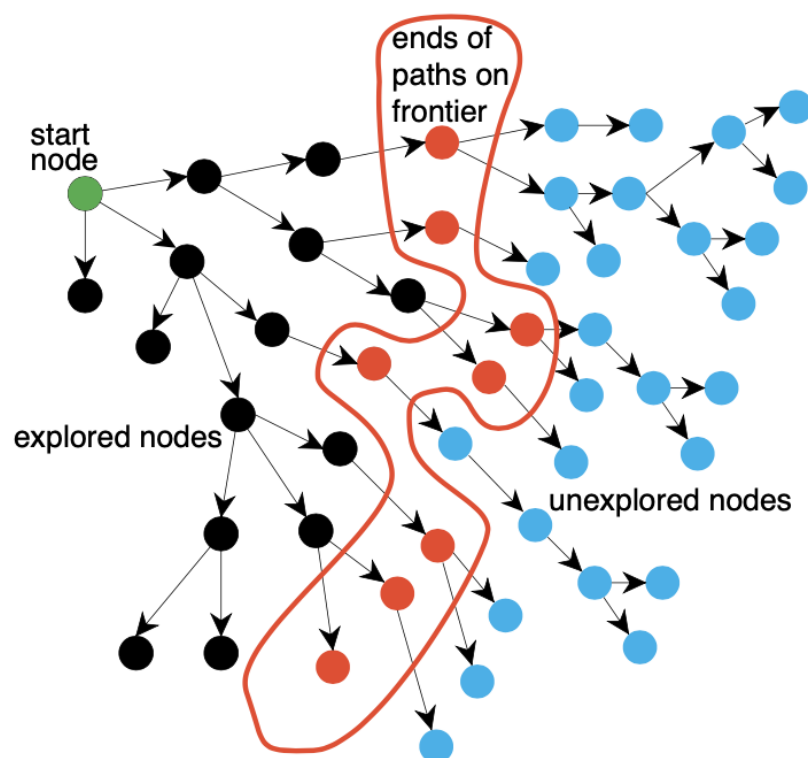


Figure 3: search tree

Generic graph search algorithm

Input: a graph,
a set of start nodes,
Boolean procedure $goal(n)$ that tests if n is a goal node
 $frontier := \{\langle s \rangle : s \text{ is a start node}\};$
while $frontier$ is not empty:
 select and remove path $\langle n_0, \dots, n_k \rangle$ from $frontier$;
 if $goal(n_k)$
 return $\langle n_0, \dots, n_k \rangle$;
 for every neighbor n of n_k
 add $\langle n_0, \dots, n_k, n \rangle$ to $frontier$;
end while

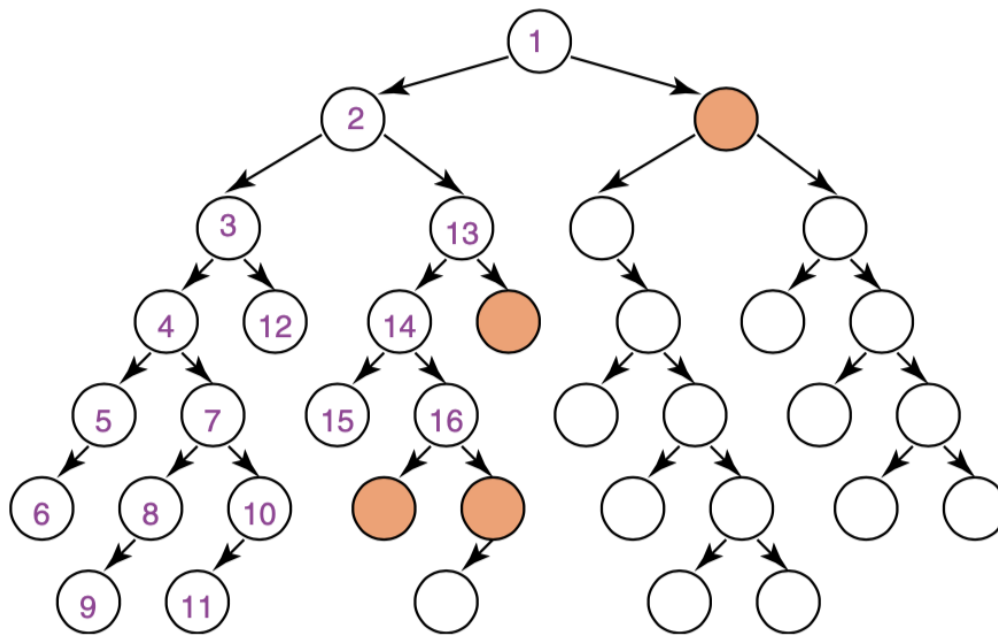
Figure 4: Generic Search

NOTE: you will have to use what ever data structure for the search you are using (BFS use a queue), (DFS use a stack).

In the generic algorithm, neighbours are going to use the method `outgoing_arcs`, we are given this algorithm in the form of a python module.

Depth-first search

- In order to perform DFS, the generic graph search must be used with a stack frontier *LIFO*
- If the stack is a python list, where each element is a path, and has the form $[\dots, p, q]$
 - q is selected and popped
 - of the algorithm continues then paths that extend q are pushed (appended) to the stack
 - p is only selected when all paths from q have been explored.
- As a result, at each stage the algorithm expands the deepest path
- The orange nodes in the graph below are considered the frontier nodes

**Figure 5:** DFS

- DFS does not guarantee a solution without pruning, due to the fact that we can have infinite loops
- It is not guaranteed to complete if it does not use pruning

A note on complexity

Assume a finite search tree of depth d and branching factor of b :

- What is the time complexity?
 - It will be exponential: $O(b^d)$
- What is the space complexity?
 - It will be linear: $O(bd)$

How do we trace the frontier

- starting with an empty frontier we record all the calls to the frontier: to add or get a path we dedicate one line per call
- When we ask the frontier to add a path, we start the line with a + followed by the path that has been added
- When we ask for a path from the frontier we start the line with a – followed by the path being removed

- When using a priority queue, the path is followed by a comma and then the key *e.g.* *cost*, *heuristic*, *f-value*, ...
- The lines of the trace should match the following regular expression $^{+ -} [a - z] + (, \backslash d +) ? ! ? \$$
- We stop when we **remove** a path from the trace

Given the following graph

```
nodes={a, b, c, d},  
edge_list=[(a,b), (a,d), (a, c), (c, d)],  
starting_nodes = [a],  
goal_nodes = {d}
```

trace the frontier in depth-first search (DFS).


Answer:

```
+ a  
- a  
+ ab  
+ ad  
+ ac  
- ac  
+ acd  
- acd
```

Figure 6: DFS trace using generic algorithm

Breath-first search

- In order to perform BFS, the generic graph search must be used with a queue frontier *FIFO*.
- If the queue is a python deque of the form $[p, q, \dots, r]$, then
 - p is selected (dequeued)
 - if the algorithm continues then paths that extend p are enqueued *appended* to the queue after r
- As a result, at each state the algorithm expands the shallowest path.



- BFS **does** guarantee to find a solution with the fewest arcs if there is a solution
- It will complete
- It will not halt due to some graphs having *cycles, with no pruning*

BFS has higher complexity than DFS

- What is the time complexity?
 - It will be exponential: $O(b^d)$
- What is the space complexity?
 - It will be linear: $O(b^d)$

Given the following graph

```
nodes={a, b, c, d},  
edge_list=[(a,b), (a,d), (a, c), (c, d)],  
starting_nodes = [a],  
goal_nodes = {d}
```

trace the frontier in breadth-first search (BFS).

Answer:

```
+ a  
- a  
+ ab  
+ ad  
+ ac  
- ab  
- ad
```

Figure 8: BFS trace using generic algorithm

Lowest-cost-first search

- The cost of a path is the sum of the costs of its arcs
- This algorithm is very similar to Dijkstra's except modified for larger graphs
- LCFS selects a path on the frontier with the lowest cost
- The frontier is a priority queue ordered by path cost
 - A priority queue is a container in which each element has a priority *cost*
 - An element with a higher priority is always selected/removed before an element with a lower priority
 - In python we can use the `heapq` you will need to store objects in a way that these properties hold
- LCFS finds an optimal solution: a least-cost path to a goal node.
- Another name for this algorithm is *uniform-cost search*.

NOTE: For an example of this queue, see Lecture One: 1:45 time stamp

Given the following graph

```
nodes={a, b, c, d, g},
edge_lists=[(a,b,4), (a,c,2), (a,d,1),
             (b,g,4), (c,g,2), (d,g,4)],
starting_nodes = [a],
goal_nodes = {g}
```

trace the frontier in lowest-cost-first search (LCFS).

Answer:

```
+ a, 0
- a, 0
+ ab, 4
+ ac, 2
+ ad, 1
- ad, 1
+ adg, 5
- ac, 2
+ acg, 4
- ab, 4
+ abg, 8
- acg, 4
```

Figure 9: LCFS trace generic

Lecture Two: Searching the State Space (part two)

Pruning

- This is our method to deal with cycles and multiple paths.
- this means we can have wasted computation and cycles in our graph

Principle: Do not expand paths to nodes that have already been expanded

Pruning Implementation

- The frontier keeps track of expanded or *closed* nodes
- When adding a new path to the frontier, it is only added if another path to the same end-node has not already been expanded, otherwise the new path is discarded (*pruned*)
- When asking for the **next path** to be returned by the frontier, a path is selected and removed but it is returned only if the end-node has not been expanded before, otherwise the path is discarded (*pruned*) and not returned. The selection and removal is repeated until a path is returned (or the frontier becomes empty). If a path is returned, its end-node will be remembered as an expanded node.

In frontier traces every time a path is pruned, we add an explanation mark ! at the end of the line

Example: LCFS with pruning

Trace LCFS with pruning on the following graph:

```
nodes = {S, A, B, G},
edge_list=[(S,A,3), (S,B,1), (B,A,1), (A,B,1), (A,G,5)],
starting_nodes = [S],
goal_nodes = {G}.
```

Answer:

```
# expanded={}
+ S,0
- S,0      # expanded={S}
+ SA,3
+ SB,1
- SB,1     # expanded={S,B}
+ SBA,2
- SBA,2    # expanded={S,B,A}
+ SBAB,3!  # not added!
+ SBAG,7
- SA,3!    # not returned!
- SBAG,7   # expanded={S,B,A,G}
```

4

Figure 10: Example: LCFS with pruning

How does LCFS behave?

- LCFS explores increasing cost contours
 - Finds an optimal solution always
 - Explores options in every direction
 - No information about goal location

We are going to use a search heuristic, function $h()$ is an estimate of the cost for the shortest path from node n to a goal node.

- h needs to be efficient to compute
- h can be extended to paths: $h(< n_0, \dots, n_k) = h(n_k)$
- h is said to be admissible if and only if:

- $\forall n \ h(n) \geq 0$, h is non-negative and $h(n) \leq C$ where C is the optimal cost of getting from n to a goal node

NOTE: We will have to come up with our own heuristic for the assignment as it depends on context.

Best-first Search

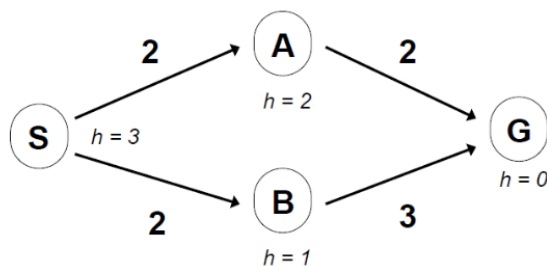
- Idea: select the path whose end is closest to a goal node according to the heuristic function.
- Best-first search is a greedy strategy that selects a path on the frontier with minimal h -value
- Main drawback: this does not guarantee finding an optimal solution.

Example: tracing best-first search

- Trace the frontier when using the best-first (greedy) search strategy for the following graph.
- The starting node is S and the goal node is G.
- Heuristic values are given next to each node.
- SA comes before SB.

heuristic function

$h(S) = 3$
 $h(A) = 2$
 $h(B) = 1$
 $h(G) = 0$



Answer:

+ S, 3
 - S, 3
 + SA, 2
 + SB, 1
 - SB, 1
 + SBG, 0
 - SBG, 0

11

Figure 11: Tracing best-first search

A search strategy

Properties:

- Always finds an optimal solution as long as:

- there is a solution
- there is no pruning
- the heuristic function is admissible
- Does it halt on every graph?

Idea:

- Don't be as wasteful as LCFS
- Don't be as greedy as best-first search
- Estimate the cost of paths as if they could be extended to reach a goal in the best possible way.

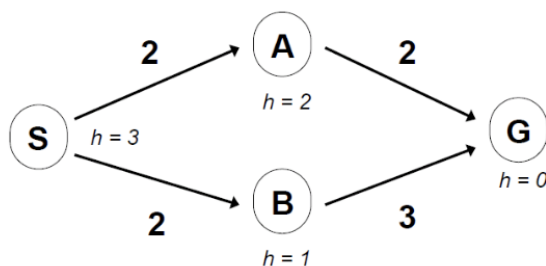
Evaluation function: $f(p) = cost(p) + h(n)$

- p is a path, n is the last node on p
- $cost(p)$ = cost of path p *this is the actual cost from the starting node to node n*
- $h(n)$ = an estimate of the cost from n to goal node
- $f(p)$ = estimated total cost of path through p to goal node

The frontier is a priority queue ordered by $f(p)$

Example: tracing A* search

- Trace the frontier when using the A* search strategy for the following graph.
- The starting node is S and the goal node is G.
- Heuristic values are given next to each node.
- SA comes before SB.



Note: This small example only show the inner working of A*. It does not demonstrate its advantage over LCFS.

heuristic function

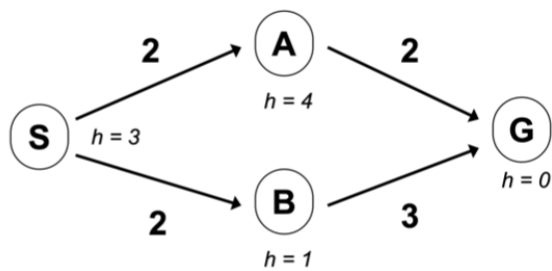
$h(S) = 3$
 $h(A) = 2$
 $h(B) = 1$
 $h(G) = 0$

Answer:

```

+ S, 3      # 0 + 3 = 3
- S, 3
+ SA, 4     # 2 + 2 = 4
+ SB, 3     # 2 + 1 = 3
- SB, 3
+ SBG, 5    # 5 + 0 = 5
- SA, 4
+ SAG, 4    # 4 + 0 = 4
- SAG, 4
  
```

- Same example as the one before just assume $h(A) = 4$ instead.



heuristic function

$$h(S) = 3$$

$$h(A) = 4$$

$$h(B) = 1$$

$$h(G) = 0$$

Answer:

$$+ S, 3 \quad \# \quad 0 + 3 = 3$$

$$- S, 3$$

$$+ SA, 6 \quad \# \quad 2 + 4 = 6$$

$$+ SB, 3 \quad \# \quad 2 + 1 = 3$$

$$- SB, 3$$

$$+ SBG, 5 \quad \# \quad 5 + 0 = 5$$

$$- SBG, 5$$

Non-optimal solution! Why?

A*: proof of optimality

When using A* (without pruning) the first path p from a starting node to a goal node that is selected and removed from the frontier has the lowest cost.

Sketch of proof:

- Suppose to the contrary that there is another path from one of the starting nodes to a goal node with a lower cost.
- There must be a path p' on the frontier such that one of its continuations leads to the goal with a lower overall cost than p .
- Since p was removed before p' :

$$f(p) \leq f(p') \implies \text{cost}(p) + h(p) \leq \text{cost}(p') + h(p') \implies \text{cost}(p) \leq \text{cost}(p') + h(p')$$

- Let c be any continuation of p' that goes to a goal node; that is, we have a path $p'c$ from a start node to a goal node. Since h is admissible, we have:

$$\text{cost}(p'c) = \text{cost}(p') + \text{cost}(c) \geq \text{cost}(p') + h(p')$$

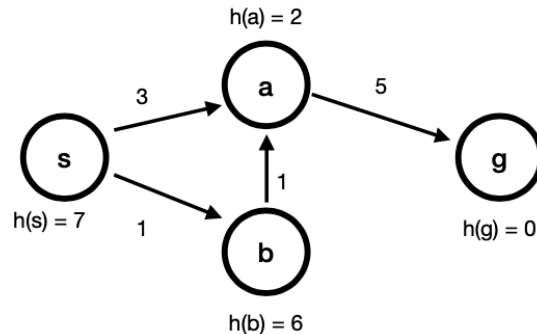
- Thus:

$$\text{cost}(p) \leq \text{cost}(p') + h(p') \leq \text{cost}(p') + \text{cost}(c) = \text{cost}(p'c)$$

Effect of pruning on A*

Trace the frontier in A* search for the following graph, with and without pruning.

```
nodes={s, a, b, g},
estimates = {s:7, a:2, b:6, g:0},
edge_list=[(s,a,3), (s,b,1),
            (b,a,1), (a,g,5)],
starting_nodes = [s],
goal_nodes = {g}.
```



Answer without pruning

```
+ S, 7
- S, 7
+ SA, 5
+ SB, 7
- SA, 5
+ SAG, 8
- SB, 7
+ SBA, 4
- SBA, 4
+ SBAG, 7
- SBAG, 7
```

Answer **with** pruning

```
# expanded={}
+ S, 7
- S, 7          # expanded={S}
+ SA, 5
+ SB, 7
- SA, 5          # expanded={S,A}
+ SAG, 8
- SB, 7          # expanded={S,A,B}
+ SBA, 4!
- SAG, 8          Non-optimal solution!
```

17

What went wrong when pruning A* Search

- An expensive path, *sa* was expanded before a cheaper path *sba* could be discovered, because $f(sa) < f(sb)$
- Is the heuristic function h admissible?
 - Yes
- So what can we do?
 - We need a stronger condition than admissibility to stop this from happening

Principle: When we are removing nodes, we are essentially saying we have found a cheaper solution, in this case, this was not true and hence why the algorithm fails, we need to use a stronger condition as outlined below

Monotonicity

A heuristic function is monotone or consistent if for every two nodes n and n' which is reachable from n :

$$h(n) \leq \text{cost}(n, n') + h(n')$$

With the monotone restriction, we have:

$$\begin{aligned} f(n') &= \text{cost}(s, n') + h(n') \\ &= \text{cost}(s, n) + \text{cost}(n, n') + h(n') \\ &\geq \text{cost}(s, n) + h(n) \\ &\geq f(n) \end{aligned}$$

How about using the actual cost as a heuristic?

- Would it be a valid heuristic?
- Would we save on nodes expanded?
- What's wrong with it?
 - It becomes as computationally expensive as it is to just do the problem

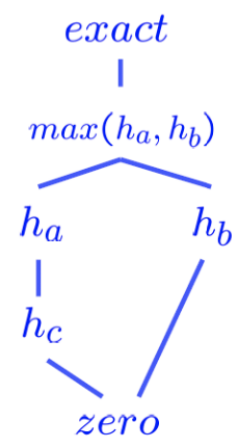
Choosing a heuristic: a trade-off between quality of estimate and work per node!

Dominance relation

- Dominance: $h_a \geq h_c$ if

$$\forall n : h_a(n) \geq h_c(n)$$
- Heuristics form a semi-lattice:
 - Max of admissible heuristics is admissible

$$h(n) = \max(h_a(n), h_b(n))$$
- Trivial heuristics
 - Bottom of lattice is the zero heuristic (what does this give us?)
 - Top of lattice is the exact heuristic



24

Figure 12: Dominance Relation