

---

## Contents

<b>Artificial Intelligence</b>	<b>2</b>
Course Information . . . . .	2
Textbooks / Resources . . . . .	2
Readings . . . . .	3
Lectures . . . . .	3
Lecture One: Searching the State Space . . . . .	3
Lecture Two: Searching the State Space (part two) . . . . .	9
Lecture Three: Knowledge Base and Information . . . . .	17
Lecture Four: Declarative Programming (Part One) . . . . .	21
Lecture Five: Declarative Programming (Part Two) . . . . .	23
Lecture Six: Local and Global Search (Optimisation) . . . . .	25
Lecture Seven: Belief Networks . . . . .	30

---

# Artificial Intelligence

## Course Information

The course covers core topics in AI including:

- uninformed and informed graph search algorithms,
- propositional logic and forward and backward chaining algorithms,
- declarative programming with Prolog,
- the min-max and alpha-beta pruning algorithms,
- Bayesian networks and probabilistic inference algorithms,
- classification learning algorithms,
- consistency algorithms,
- local search and heuristic algorithms such as simulated annealing, and population-based algorithms such as genetic search and swarm optimisation.

## Grades

Standard Computer science policy applies

- Average 50% over all assessment items
- Average at least 45% on all invigilated assessment items

Grading structure for course

- Assignments (5%)
  - Two Super Quiz's
- Quizzes (16.5%)
  - Weekly Quiz Assessments (1.5% ea)
- Lab Test (20%)
- Final Exam (58.5%)

## Textbooks / Resources

- Poole, David L. 1958, Mackworth, Alan K; Artificial intelligence : foundations of computational agents; Cambridge University Press, 2010.
- Russell, Stuart J, Norvig, Peter; Artificial intelligence : a modern approach; 3rd ed; Prentice Hall, 2010.

---

## Readings

## Lectures

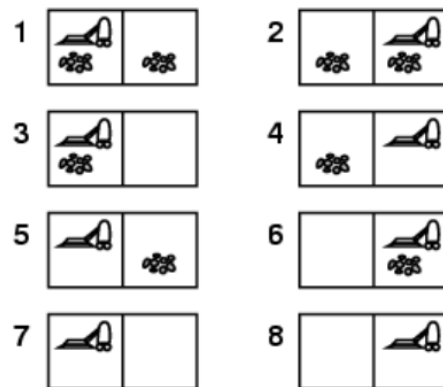
### Lecture One: Searching the State Space

#### What is state?

- A state is a data structure that represents a possible configuration of the world *agent and environment*
  - The **state space** is the set of all possible states for that problem
  - actions change the state of the world
- 
- Example: A vacuum cleaner agent in two adjacent rooms which can be either clean or dirty.

- Location = {left, right}
- Left-room-condition = {dirty, clean}
- Right-room-condition = {dirty, clean}
- State-space = Location  
× Left-room-condition  
× Right-room-condition

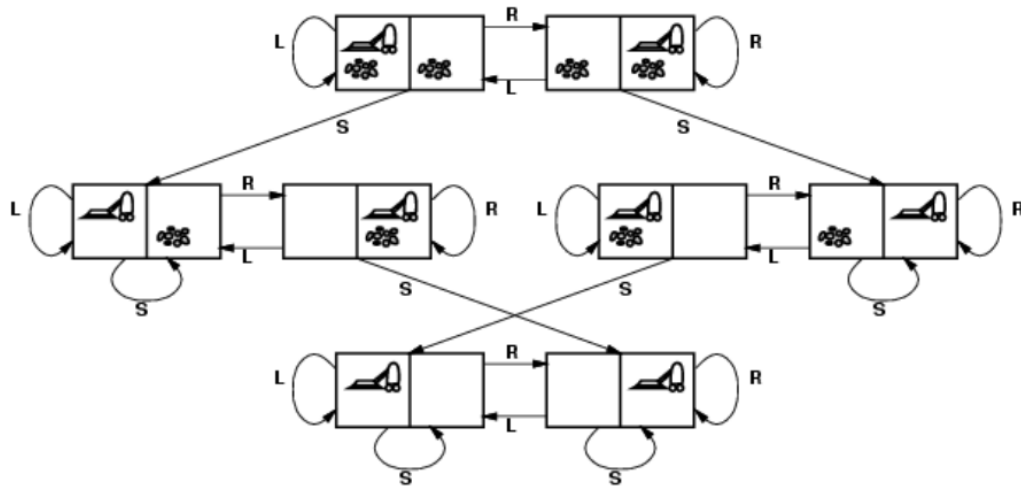
In this example, each state is represented by a triple (3-tuple).



**Figure 1:** State space example one

State can also be represented as a graph *both directed and undirected*

- Example: Suppose the vacuum cleaner agent can take the following actions: L (go left), R (go right), S (suck).



**Figure 2:** State space graph simplified

- Many problems in AI can be abstracted to the problem of finding a path in a directed graph
- Notation we use is **Nodes** and **arcs** for **vertices** and **edges** in a graph

### Explicit vs Implicit graphs

- In **explicit graphs** nodes and arcs are readily available, they are read from the input and stored in a data structure such as an adjacency list/matrix.
  - the entire graph is in memory.
  - the complexity of algorithms are measured in the number of nodes and/or arcs.
- In **implicit graphs** a procedure `outgoing_arcs` is defined that given a node, returns a set of directed arcs that connect node to other nodes.
  - The graph is generated as needed *due to the complexity of the graphs*.
  - The complexity is measured in terms of the depth of the goal state node or *how far do we have to get into the graph to find a solution*.

### Explicit graphs in quizzes

- In some exercises we use small explicit graphs to study the behaviour of various frontiers
- Nodes are specified in a set

- 
- Edges are specified in a list
    - pairs of nodes, or triples of nodes (in a tuple)

### Searching graphs

- We will use generic search algorithms: given a graph, start nodes, and goal nodes, incrementally explore paths from the start nodes.
- Maintain a **frontier** of paths that have been explored
  - frontier: paths that we have already explored
- As search proceeds, the frontier is updated and the graph is explored until a goal node is found.
- The order in which paths are removed and added to the frontier defines the search strategy
- A **search tree** is a tree drawn out of all the possible actions in terms of a tree.
  - How do we handle loops? *Covered in next lecture*
  - In the search tree outlined below, you can see that the *end of paths on frontier* represents a BFS relationship note this is not always the case.

search tree

### Generic graph search algorithm

---

```

Input: a graph,
        a set of start nodes,
        Boolean procedure goal(n) that tests if n is a goal node
frontier := {⟨s⟩ : s is a start node};
while frontier is not empty:
    select and remove path ⟨n0, ..., nk⟩ from frontier;
    if goal(nk)
        return ⟨n0, ..., nk⟩;
    for every neighbor n of nk
        add ⟨n0, ..., nk, n⟩ to frontier;
end while

```

**Figure 3:** Generic Search

NOTE: you will have to use what ever data structure for the search you are using (BFS use a queue), (DFS use a stack).

In the generic algorithm, neighbours are going to use the method `outgoing_arcs`, we are given this algorithm in the form of a python module.

### Depth-first search

- In order to perform DFS, the generic graph search must be used with a stack frontier *LIFO*
- If the stack is a python list, where each element is a path, and has the form [..., p, q]
  - *q* is selected and popped
  - of the algorithm continues then paths that extend *q* are pushed (appended) to the stack
  - *p* is only selected when all paths from *q* have been explored.
- As a result, at each stage the algorithm expands the deepest path
- The orange nodes in the graph below are considered the frontier nodes

DFS

- 
- DFS does not guarantee a solution without pruning, due to the fact that we can have infinite loops
  - It is not guaranteed to complete if it does not use pruning

### A note on complexity

Assume a finite search tree of depth  $d$  and branching factor of  $b$ :

- What is the time complexity?
  - It will be exponential:  $O(b^d)$
- What is the space complexity?
  - It will be linear:  $O(bd)$

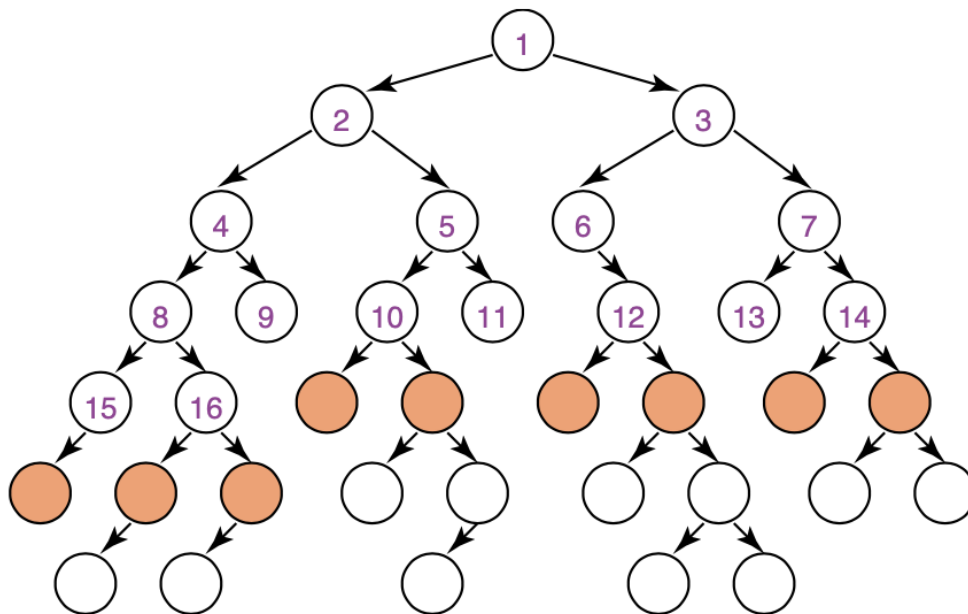
### How do we trace the frontier

- starting with an empty frontier we record all the calls to the frontier: to add or get a path we dedicate one line per call
- When we ask the frontier to add a path, we start the line with a + followed by the path that has been added
- When we ask for a path from the frontier we start the line with a – followed by the path being removed
- When using a priority queue, the path is followed by a comma and then the key *e.g. cost, heuristic, f-value, ...*
- The lines of the trace should match the following regular expression  $^{+ -} [a-z] + (, \backslash d +) ? ! ? \$$
- We stop when we **remove** a path from the trace

DFS trace using generic algorithm

### Breadth-first search

- In order to perform BFS, the generic graph search must be used with a queue frontier *FIFO*.
- If the queue is a python deque of the form  $[p, q, \dots, r]$ , then
  - $p$  is selected (dequeued)
  - if the algorithm continues then paths that extend  $p$  are enqueued *appended* to the queue after  $r$
- As a result, at each state the algorithm expands the shallowest path.



**Figure 4:** BFS Illustration of search tree

- BFS **does** guarantee to find a solution with the fewest arcs if there is a solution
- It will complete
- It will not halt due to some graphs having cycles, *with no pruning*

#### A note on complexity

BFS has higher complexity than DFS

- What is the time complexity?
  - It will be exponential:  $O(b^d)$
- What is the space complexity?
  - It will be linear:  $O(b^d)$

BFS trace using generic algorithm

#### Lowest-cost-first search

- The cost of a path is the sum of the costs of its arcs
- This algorithm is very similar to Dijkstra's except modified for larger graphs
- LCFS selects a path on the frontier with the lowest cost
- The frontier is a priority queue ordered by path cost



- 
- A priority queue is a container in which each element has a priority *cost*
  - An element with a higher priority is always selected/removed before an element with a lower priority
  - In python we can use the `heapq` you will need to store objects in a way that these properties hold
- LCFS finds an optimal solution: a least-cost path to a goal node.
  - Another name for this algorithm is *uniform-cost search*.

NOTE: For an example of this queue, see Lecture One: 1:45 time stamp

LCFS trace generic

## Lecture Two: Searching the State Space (part two)

### Pruning

- This is our method to deal with cycles and multiple paths.
- this means we can have wasted computation and cycles in our graph

Principle: Do not expand paths to nodes that have already been expanded

### Pruning Implementation

- The frontier keeps track of expanded or *closed* nodes
- When adding a new path to the frontier, it is only added if another path to the same end-node has not already been expanded, otherwise the new path is discarded (*pruned*)
- When asking for the **next path** to be returned by the frontier, a path is selected and removed but it is returned only if the end-node has not been expanded before, otherwise the path is discarded (*pruned*) and not returned. The selection and removal is repeated until a path is returned (or the frontier becomes empty). If a path is returned, its end-node will be remembered as an expanded node.

In frontier traces every time a path is pruned, we add an explanation mark ! at the end of the line

---

# Example: LCFS with pruning

Trace LCFS with pruning on the following graph:

```
nodes = {S, A, B, G},  
edge_list=[(S,A,3), (S,B,1), (B,A,1), (A,B,1), (A,G,5)],  
starting_nodes = [S],  
goal_nodes = {G}.
```

Answer:

```
# expanded={}  
+ S,0  
- S,0      # expanded={S}  
+ SA,3  
+ SB,1  
- SB,1     # expanded={S,B}  
+ SBA,2  
- SBA,2    # expanded={S,B,A}  
+ SBAB,3!  # not added!  
+ SBAG,7  
- SA,3!    # not returned!  
- SBAG,7   # expanded={S,B,A,G}  
  
4
```

**Figure 5:** Example: LCFS with pruning

## How does LCFS behave?

- LCFS explores increasing cost contours
  - Finds an optimal solution always
  - Explores options in every direction
  - No information about goal location

We are going to use a search heuristic, function  $h()$  is an estimate of the cost for the shortest path from node  $n$  to a goal node.

- $h$  needs to be efficient to compute
- $h$  can be extended to paths:  $h(< n_0, \dots, n_k) = h(n_k)$
- $h$  is said to be admissible if and only if:
  - $\forall n \ h(n) \geq 0$ ,  $h$  is non-negative and  $h(n) \leq C$  where  $C$  is the optimal cost of getting from  $n$  to a goal node

---

NOTE: We will have to come up with our own heuristic for the assignment as it depends on context.

### Best-first Search

- Idea: select the path whose end is closest to a goal node according to the heuristic function.
- Best-first search is a greedy strategy that selects a path on the frontier with minimal  $h$ -value
- Main drawback: this does not guarantee finding an optimal solution.

## Example: tracing best-first search

- Trace the frontier when using the best-first (greedy) search strategy for the following graph.
- The starting node is S and the goal node is G.
- Heuristic values are given next to each node.
- SA comes before SB.

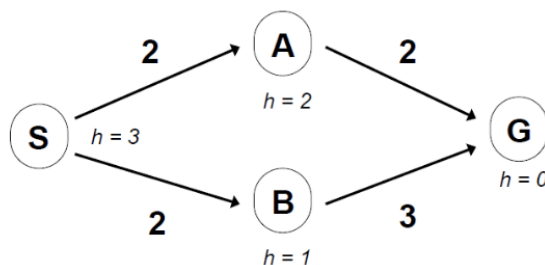
heuristic function

$$h(S) = 3$$

$$h(A) = 2$$

$$h(B) = 1$$

$$h(G) = 0$$



Answer:

+ S, 3

- S, 3

+ SA, 2

+ SB, 1

- SB, 1

+ SBG, 0

- SBG, 0

11

**Figure 6:** Tracing best-first search

### A search strategy

Properties:

- Always finds an optimal solution as long as:
  - there is a solution
  - there is no pruning

- the heuristic function is admissible

- Does it halt on every graph?

Idea:

- Don't be as wasteful as LCFS
- Don't be as greedy as best-first search
- Estimate the cost of paths as if they could be extended to reach a goal in the best possible way.

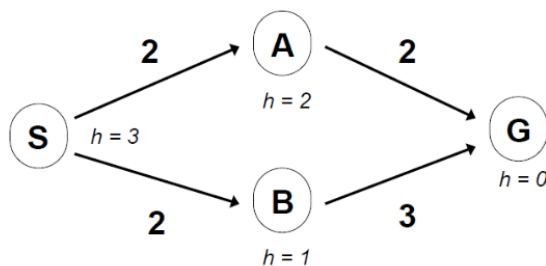
Evaluation function:  $f(p) = cost(p) + h(n)$

- $p$  is a path,  $n$  is the last node on  $p$
- $cost(p)$  = cost of path  $p$  this is the actual cost from the starting node to node  $n$
- $h(n)$  = an estimate of the cost from  $n$  to goal node
- $f(p)$  = estimated total cost of path through  $p$  to goal node

The frontier is a priority queue ordered by  $f(p)$

## Example: tracing A\* search

- Trace the frontier when using the A\* search strategy for the following graph.
- The starting node is S and the goal node is G.
- Heuristic values are given next to each node.
- SA comes before SB.



**Note:** This small example only show the inner working of A\*. It does not demonstrate its advantage over LCFS.

heuristic function

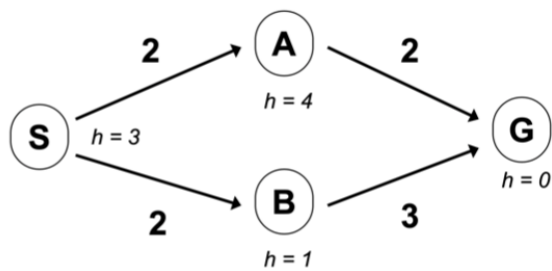
$h(S) = 3$   
 $h(A) = 2$   
 $h(B) = 1$   
 $h(G) = 0$

Answer:

```

+ S, 3      # 0 + 3 = 3
- S, 3
+ SA, 4     # 2 + 2 = 4
+ SB, 3     # 2 + 1 = 3
- SB, 3
+ SBG, 5    # 5 + 0 = 5
- SA, 4
+ SAG, 4    # 4 + 0 = 4
- SAG, 4
  
```

- Same example as the one before just assume  $h(A) = 4$  instead.



heuristic function

$$h(S) = 3$$

$$h(A) = 4$$

$$h(B) = 1$$

$$h(G) = 0$$

Answer:

$$+ S, 3 \quad \# \quad 0 + 3 = 3$$

$$- S, 3$$

$$+ SA, 6 \quad \# \quad 2 + 4 = 6$$

$$+ SB, 3 \quad \# \quad 2 + 1 = 3$$

$$- SB, 3$$

$$+ SBG, 5 \quad \# \quad 5 + 0 = 5$$

$$- SBG, 5$$

Non-optimal solution! Why?

---

# A\*: proof of optimality

When using A\* (without pruning) the first path  $p$  from a starting node to a goal node that is selected and removed from the frontier has the lowest cost.

Sketch of proof:

- Suppose to the contrary that there is another path from one of the starting nodes to a goal node with a lower cost.
- There must be a path  $p'$  on the frontier such that one of its continuations leads to the goal with a lower overall cost than  $p$ .
- Since  $p$  was removed before  $p'$ :

$$f(p) \leq f(p') \implies \text{cost}(p) + h(p) \leq \text{cost}(p') + h(p') \implies \text{cost}(p) \leq \text{cost}(p') + h(p')$$

- Let  $c$  be any continuation of  $p'$  that goes to a goal node; that is, we have a path  $p'c$  from a start node to a goal node. Since  $h$  is admissible, we have:

$$\text{cost}(p'c) = \text{cost}(p') + \text{cost}(c) \geq \text{cost}(p') + h(p')$$

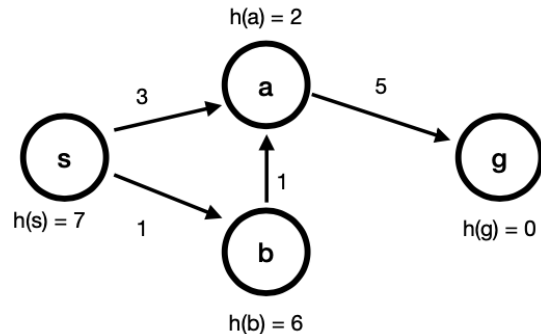
- Thus:

$$\text{cost}(p) \leq \text{cost}(p') + h(p') \leq \text{cost}(p') + \text{cost}(c) = \text{cost}(p'c)$$

# Effect of pruning on A\*

Trace the frontier in A\* search for the following graph, with and without pruning.

```
nodes={s, a, b, g},
estimates = {s:7, a:2, b:6, g:0},
edge_list=[(s,a,3), (s,b,1),
            (b,a,1), (a,g,5)],
starting_nodes = [s],
goal_nodes = {g}.
```



## Answer without pruning

```
+ S, 7
- S, 7
+ SA, 5
+ SB, 7
- SA, 5
+ SAG, 8
- SB, 7
+ SBA, 4
- SBA, 4
+ SBAG, 7
- SBAG, 7
```

## Answer **with** pruning

```
# expanded={}
+ S, 7
- S, 7          # expanded={S}
+ SA, 5
+ SB, 7
- SA, 5          # expanded={S,A}
+ SAG, 8
- SB, 7          # expanded={S,A,B}
+ SBA, 4!
- SAG, 8          Non-optimal solution!
```

17

## What went wrong when pruning A Search

- An expensive path, *sa* was expanded before a cheaper path *sba* could be discovered, because  $f(sa) < f(sb)$
- Is the heuristic function  $h$  admissible?
  - Yes
- So what can we do?
  - We need a stronger condition than admissibility to stop this from happening

Principle: When we are removing nodes, we are essentially saying we have found a cheaper solution, in this case, this was not true and hence why the algorithm fails, we need to use a stronger condition as outlined below

## Monotonicity

A heuristic function is monotone or consistent if for every two nodes  $n$  and  $n'$  which is reachable from  $n$ :

---

$$h(n) \leq \text{cost}(n, n') + h(n')$$

With the monotone restriction, we have:

$$\begin{aligned} f(n') &= \text{cost}(s, n') + h(n') \\ &= \text{cost}(s, n) + \text{cost}(n, n') + h(n') \\ &\geq \text{cost}(s, n) + h(n) \\ &\geq f(n) \end{aligned}$$

How about using the actual cost as a heuristic?

- Would it be a valid heuristic?
- Would we save on nodes expanded?
- What's wrong with it?
  - It becomes as computationally expensive as it is to just do the problem

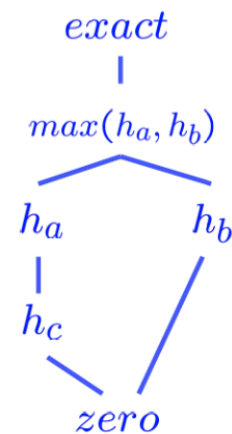
Choosing a heuristic: a trade-off between quality of estimate and work per node!



---

# Dominance relation

- Dominance:  $h_a \geq h_c$  if
$$\forall n : h_a(n) \geq h_c(n)$$
- Heuristics form a semi-lattice:
  - Max of admissible heuristics is admissible
$$h(n) = \max(h_a(n), h_b(n))$$
- Trivial heuristics
  - Bottom of lattice is the zero heuristic (what does this give us?)
  - Top of lattice is the exact heuristic



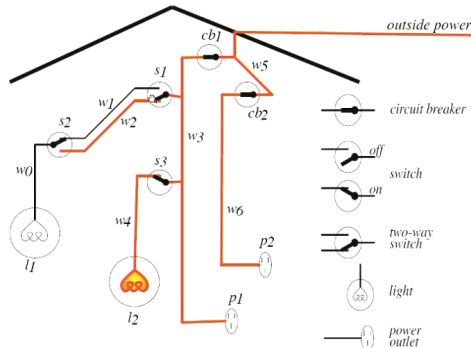
24

**Figure 7:** Dominance Relation

Further algorithms are discussed in this segment of the and lecture *boards onto lecture three* however this content will not be assessed in the duration of this course.

## Lecture Three: Knowledge Base and Information

### How to represent information in a knowledge base



## Representing the Electrical Environment

```

light.j1.      lit.j1 ← live.w0 ∧ ok.j1
live.w0 ← live.w1 ∧ up.s2.
light.j2.      live.w0 ← live.w2 ∧ down.s2.
down.s1.       live.w1 ← live.w3 ∧ up.s1.
up.s2.         live.w2 ← live.w4 ∧ down.s1.
up.s3.         lit.j2 ← live.w4 ∧ ok.j2.
ok.j1.         live.w4 ← live.w3 ∧ up.s3.
ok.j2.         live.p1 ← live.w3.
ok.cb1.        live.w3 ← live.w5 ∧ ok.cb1.
ok.cb2.        live.p2 ← live.w6.
live.outside.  live.w5 ← live.w5 ∧ ok.cb2.
live.w5 ← live.outside.

```

Chapter 5 Propositions and Inference 3/22

- The computer doesn't know the meaning of the symbols (logical and etc...)
- The user can interpret the symbol using their meaning
- There is no specific syntax for this, it is just what ever is readable for the user/writer

## Simple language and definitions

- An **atom** is a symbol starting with a lower case letter
- A **body** is an atom or is of the form  $b_1 \wedge b_2$  where  $b_1$  and  $b_2$  are bodies
- A **definite clause** is an atom or a rule of the form  $h \leftarrow b$  where  $h$  is an atom and  $b$  is a body
- A **knowledge base** is a set of definite clauses
- An **interpretation**  $i$  assigns a truth value to each atom
- A **body**  $b_1 \wedge b_2$  is true in  $i$  if  $b_1$  is true in  $i$  and  $b_2$  is true in  $i$
- A **rule**  $h \leftarrow b$  is false in  $i$  if  $b$  is true in  $i$  and  $h$  is false in  $i$ , the rule is true otherwise
- A **knowledge base**  $KB$  is true in  $i$  if and only if every clause in  $KB$  is true in  $i$
- A **model** of a set of clauses is an interpretation in which all the clauses are *true*
- If  $KB$  is a set of clauses and  $g$  is a conjunction of atoms,  $g$  is a **logical consequence** of  $KB$ , this is denoted as  $KB \models g$ , if  $g$  is *true* in every model of  $KB$ 
  - That is,  $KB \models g$  if there is no interpretation in which  $KB$  is *true* and  $g$  is *false*.
- A **Proof procedure** is a -possibly non-deterministic - algorithm for deriving consequences of a knowledge
- Given a proof procedure,  $KB \vdash g$  means  $g$  can be derived from knowledge base  $KB$
- Recall  $KB \models g$  means  $g$  is *true* in all models of  $KB$
- A proof procedure is **sound** if  $KB \vdash g \implies KB \models g$
- A proof procedure is **complete** if  $KB \models g \implies KB \vdash g$

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

How many interpretations?

	$p$	$q$	$r$	$s$	model of $KB$ ?
$I_1$	true	true	true	true	
$I_2$	false	false	false	false	
$I_3$	true	true	false	false	
$I_4$	true	true	true	false	
$I_5$	true	true	false	true	

Which of  $p, q, r, s$  logically follow from  $KB$ ?

**Figure 8:** simple example question

Answers to the questions:

We have four atoms  $\{p, q, r, s\}$ , because we have 4 atoms, there are 16 permutations in our truth table ( $2^4$ ), therefore we have 16 interpretations

### Bottom-up proof procedure

Rule of derivation:

if  $h \leftarrow b_1 \wedge \dots \wedge \dots b_m$  is a clause in the knowledge base, and each  $b_i$  has been derived, then  $h$  can be derived

- This is **Forward chaining** on this clause (this rule also covers the case when  $m = 0$ )
- $KB \vdash g$  if  $g \in C$  at the end of the below algorithmic procedure

- Tracing tutorial: 1:13:30

```

C := {};
repeat
    select clause " $h \leftarrow b_1 \wedge \dots \wedge b_m$ " in  $KB$  such that
         $b_i \in C$  for all  $i$ , and
         $h \notin C$ ;
    C := C  $\cup$  {h}
until no more clauses can be selected.

```

**Figure 9:** Bottom-up proof procedure algorithm pseudo code

### Top-down proof procedure

Idea: search backward from a query to determine if it is a logical consequence of  $KB$

An **answer clause** is of the form:

- $yes \leftarrow a_i \wedge \dots \wedge a_m$

The SLD Resolution of this answer clause on atom  $a_i$  with the clause:

- $a_i \leftarrow b_1 \wedge \dots \wedge b_p$
- Tracing tutorial: 1:31:00

An **answer** is an answer clause with  $m = 0$ . That is the answer clause  $yes \leftarrow$ .

A **Derivation** of query  $?q_1 \wedge \dots \wedge q_k$  from  $KB$  is a sequence of answer clauses  $\lambda_0, \lambda_1, \dots, \lambda_n$

- $\lambda_0$  is the answer clause  $yes \leftarrow q_1 \wedge \dots \wedge q_k$
- $\lambda_1$  is obtained by resolving  $\lambda_{i-1}$  with a clause in  $KB$
- $\lambda_n$  is the answer

---

To solve the query  $?q_1 \wedge \dots \wedge q_k$ :

$ac := \text{"yes"} \leftarrow q_1 \wedge \dots \wedge q_k$

**repeat**

**select** atom  $a_i$  from the body of  $ac$

**choose** clause  $C$  from  $KB$  with  $a_i$  as head

    replace  $a_i$  in the body of  $ac$  by the body of  $C$

**until**  $ac$  is an answer.

**Figure 10:** Top-down proof procedure algorithm pseudo code

There is more information on SLD Resolution at the end of this lecture, this will be needed in the assignment

## Lecture Four: Declarative Programming (Part One)

### What is declarative programming?

Declarative programming is the use of mathematical logic to describe the logic of computation without describing its control flow

- Knowledge bases and queries in propositional logic are made up of propositions and connectives
- Predicate logic adds the notion of *predicates* and *variables*
- We take a non-theoretical approach to predicate logic by introducing *declarative programming*
- useful for: expert systems, diagnostics, machine learning, parsing text, theorem proving, ...

### Datalog

- Prolog is a declarative programming language and stand for PROgramming in LOGic
- we only look at a subset of the language which is equal to Datalog
- Think declaratively, not procedurally
- High level, interpreted language
- We will have a file that contains a knowledge base, and we will have an interpreter where we can ask queries

Here is an example of a knowledge base in Datalog:

---

```
1 woman(mia)
2 woman(jody)
3 woman(yolanda)
4 playsAirGuitar(yolanda)
```

Here is how we may query data using the interpreter:

```
1 $ woman(mia)
2 yes
```

Further examples of this are in the slides of lecture four

## Operators

- Implication :-
- Conjunction: , (AND)
- Disjunction ; (OR)
- We will later talk about how to simulate the (NOT) operator

Interpreter Operands and rules:

- Variables: X, Y, Z, Cam, AnythingThatStartswithUppercase
  - Acts as a **wildcard** to match with when querying
- Order of arguments matters
- **Arity** is important
- Unification/matching:
  - Two terms unify or match if they are the same term or if they contain variables that can be uniformly instantiated with terms in such a way that the resulting terms are equal (this is how we query)
  - Example:  $l(s(g), Z) = k(X, t(Y))$

With only Unification we can do some programming

```
1 vertical(line(point(X,Y), point(X,Z)))
2 horizontal(line(point(X,Y), point(Z,Y)))
```

## Proof Search

- Prolog has a specific way of answering queries
  - Search knowledge base from top to bottom
  - Processes clauses from left to right

- 
- Backtracking to recover from bad choices
  - Further examples using prolog: 1:10:00

### Recursive Programming

```
1 child(anna, bridget)
2 child(bridget, caroline)
3 child(caroline, donna)
4 child(donna, emily)
5 decendent(X,Y):-child(X,Y)
6 decendent(X,Y):-child(X,Z), decendent(Z,Y)
```

If we make the following query with the above knowledge base, we get a positive response

```
1 $- decendent(anna, donna)
2 yes
```

## Lecture Five: Declarative Programming (Part Two)

### Lists in Prolog

- A list is a finite sequence of elements
- List elements are enclosed in square brackets
- we can think of non-empty lists as a head and tail
  - Head is first item
  - Tail is the rest of the list
- Empty list has no head or tail
- Here are some examples of lists in prolog

```
1 [mia, vincent, jules, yolanda]
2 [mia, robber(honeybunny), X, 2, mia]
3 []
```

### Pipe Operand

- Can be used for creating a list
- Example:

```
1 [head|tail] = [mia, vincent, jules, yolanda].
2 Head = mia
3 tail = [vincent, jules, yolanda].
```

- We can have anonymous variables denoted with the \_

- 
- These do not get recorded and assigned to variables

```
1 [_ ,X2,_,X4|_] = [mia, vincent, jules, yolanda].
2 X2 = vincent
3 X4 = Jody
```

### Defining Members of a list

- One of the most basic things we would like to know is whether something is an element of a list or not
- So let's write a predicate that when given a term  $X$  and a list  $L$ , tells us whether  $X \in L$
- We can define member as the following:

```
1 member(X,[X,_]).
2 member(X,_,T):-member(X,T).
```

### Defining Append

- We can define an important predicate, append whose arguments are all lists
- Declaratively, append(L1,L2,L3) is true if list L3 is the result of concat L1, L2
- Recursive definition,
  - Base case: appending the empty list to any list produces the same list
  - The recursive step says that when concatenating non-empty list  $[H|T]$  with list  $L$ , the result is a list with head  $H$  and the result of concatenating  $T$  and  $L$

Definition:

```
1 append([],L,L).
2 append([H|L1],L2,[H|L3]):-append(L1,L2,L3).
```

Expected Output:

```
1 $- append([a,b,c],[d,e,f], Z).
2 $- Z = [a,b,c,d,e,f].
3 yes
```

### Sublist

- Now it is very easy to write a predicate that finds sub-lists of lists
- The sub-lists of a list  $L$  are simply the prefixes or suffixes of  $L$
- Checks if a list is a subset of another list

```
1 sublist(Sub,List):-suffix(Suffix,List),prefix(Sub,Suffix).
```

### Reversing a list



- 
- Recursive definition

1. If we reverse the empty list, we obtain the empty list
2. If we reverse the list  $[H|T]$ , we end up with the list obtained by reversing  $T$
3. This solution works, but is extremely inefficient, *Quadratic time*

```
1 reverse([], []).
2 reverse([H|T], R) :- reverse(T, RT), append(RT, [H], R).
```

- Here is a much more efficient solution:
- We can use an accumulator (list to append the reverse to) in order to make this faster

```
1 accReverse([], L, L).
2 accReverse([H|T], Acc, Rev) :- accReverse(T, [H|Acc], Rev).
3
4 reverse(L1, L2) :- accReverse(L1, [], L2). # Wrapper for accReverse function
```

The above is a more efficient solution

### Negation as Failure

- We need to use the cut operator (!) to suppress backtracking
- The fail predicate always fails
- They can be combined to get a negation as failure

```
1 neg(Goal) :- Goal, !, fail.
2 neg(Goal).
```

## Lecture Six: Local and Global Search (Optimisation)

### Optimisation Problems

Given:

- A set of variables and their domains; and
- An objective function (aka a cost function),

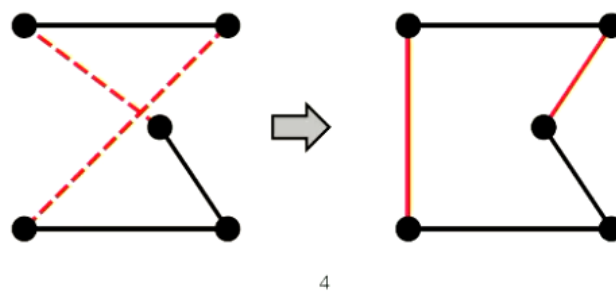
Find an assignment (of each value to each variable) that optimises (max or min) the value of the objective function.

- Optimisation usually involves searching
- CSP's and optimisation problems can be converted (reduced) to each other
- There are special algorithms for certain kind of optimisation problems (linear programming, convex optimisation)

- 
- In this lecture we will look at two families of algorithms (local and global)

### Local Search for Optimisation

- A **Local search** algorithm is an iterative algorithm that keeps a single current state and in each iteration tries to improve it by moving to one of its neighbouring states.
- Two key aspects to decide:
  - Neighbourhood: which states are the neighbours of a given state
  - Movement: which neighbouring state should the algorithm go to
- A search algorithm is considered to be greedy if it always moves to the best neighbour. Two variants happen to have special names:
  - *Hill climbing*: for maximisation
  - *Greedy descent*: for minimisation
- Traveling Salesperson Problem (TSP): Given a list of cities and the distances between each pair of cities, what is the shortest possible route that visits each city exactly once and returns to the origin city?
- Start with any complete tour, in each iteration perform pairwise exchanges if it improves the total cost.
- Variants of this approach can get close to optimal solution quickly (even with a large number of cities).



**Figure 11:** Example of local search

Local search for CSP's:

- A constrained satisfaction problem can be reduced to an optimisation problem

- 
- Given an assignment, a conflict is an unsatisfied constraint
  - Heuristic function: the number of conflicts produced by an assignment
  - Optimisation problem: find an assignment that minimises this heuristic function

Local search for CSP's in neighbourhood:

- Neighbours of a given state can be defined in many ways
  - All possible assignments except the current one
  - Select a variable that appears in any conflict, neighbours are assignment in which that variable takes a different value from its domain
  - Select a variables in the current assignment that participates in the most number of conflicts. Neighbours are assignments in which that variable takes a different value from its domain
  - n-queens example (35:00)

## Global Search

Parallel search:

- A total assignment is called an **individual**
- Idea: maintain a population of  $k$  individuals instead of one
- At every stage, update each individual in the population
- Like  $k$  restarts, but uses  $k$  times the minimum number of steps
- A basic form of global search

Simulating Annealing:

- Pick a variable at random and a new value at random
- If it is an improvement, adopt it
- If it isn't an improvement adopt it probabilistically depending on a temperature parameter,  $T$
- Temperature can be reduced

Temperature	1-worse	2-worse	3-worse
10	0.91	0.81	0.74
1	0.37	0.14	0.05
0.25	0.02	0.003	0.00005
0.1	0.00005	0	0

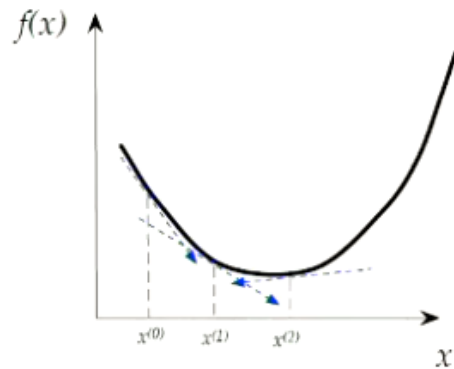
## Gradient Descent

- A widely used local search algorithm in numeric optimisation (machine learning)
- Used when the variables are numeric and continuous
- The objective function must be differentiable (mostly)

```

1: Guess  $\mathbf{x}^{(0)}$ , set  $k \leftarrow 0$ 
2: while  $\|\nabla f(\mathbf{x}^{(k)})\| \geq \epsilon$  do
3:    $\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - t_k \nabla f(\mathbf{x}^{(k)})$ 
4:    $k \leftarrow k + 1$ 
5: end while
6: return  $\mathbf{x}^{(k)}$ 

```



**Figure 12:** Gradient descent algorithm

## Genetic Algorithms

- Genetic algorithms and the whole family of evolutionary algorithms are inspired by natural selection
- They are in the global search family
- The algorithm maintains a population of individuals which evolves over time
- We can make an individual to represent anything we want
- A fitness function is needed. The function takes an individual as input and returns a numeric number indicating how good/bad the individual is
- A mechanism is needed to create the initial population
- A mechanism is needed to **evolve** the current population to the next one. This involves the following mechanisms:
  - Selection: decide which individuals survive or can reproduce
  - Crossover: given a number of parent individuals, create a number of children
  - Mutation: make some random changes to individuals
  - How this works (1:20:00)

## Roulette Wheel selection: Example

- Sum the fitness of all individuals, call it  $T$
- Generate a random number  $N$  between 1 and  $T$
- Return individual whose fitness added to the running total is equal to or larger than  $N$

- Chance to be selected is exactly proportional to fitness
- Individual: [1,2,3,4,5,6]
- Fitness: [8,2,17,7,4,11]
- Running total: [8,10,27,34,38,49]
- $N$  : 23
- selected: 3

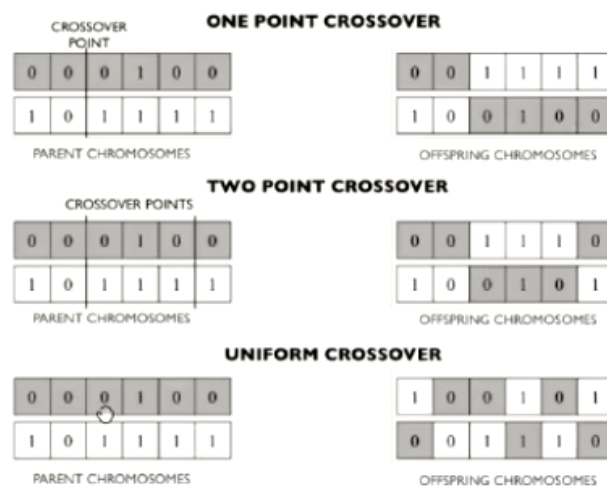
### Tournament Selection

- Choose  $n$  individuals randomly; the fittest one is selected as a parent
- $n$  is the **size** of the tournament
- By changing the size, selection pressure can be adjusted

### Crossover

Often individuals are represented as a sequence (tuple) of values. With this representation, cross over can be performed very easily.

- Generate 1,2, or a number of random *crossover points*
- Split the parents at these points
- Create offspring's by exchanging alternative segments



**Figure 13:** Crossover Example

### Mutation

With sequential representation (tuples), mutation is performed by selecting one or more random locations (indices) and changing the values at those locations to some random values (from the domain).

---

## Mutation vs Crossover

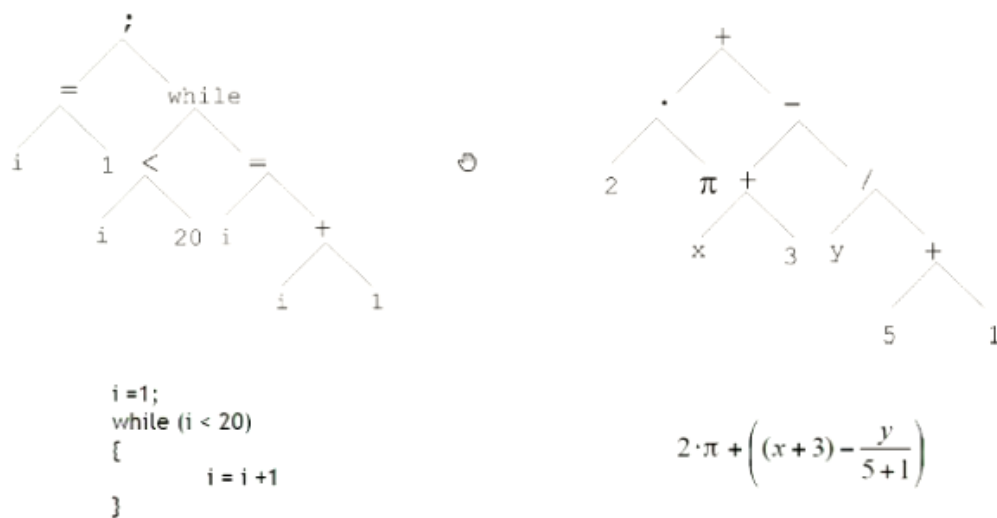
- Purpose of crossover: combining somewhat good candidates in the hope of producing better children
- Purpose of mutation: bringing diversity
- Its good to have both
- Mutation-only-EA is possible, crossover-only-EA would not work

## Fitness landscapes

- EA's are known to be able to handle relatively challenging fitness landscapes
- *see lecture: sep 14 for more information, note to self: go over this again before final*

## Tree representation

- Individuals can have more sophisticated structure
- The following shows two example trees representing statements and expressions



**Figure 14:** Tree representation example

## Lecture Seven: Belief Networks

What are belief networks about?

- Long answer short **Probabilities**

- 
- Reasons for uncertainty and randomness

### Random Variables

- A random variable is some aspect of the world which we have uncertainty
  - R = Is it raining?
  - D = How long will it take to drive to work?
  - L = Where am I?
- We denote random variables with capital letters
- Each random variable has a domain
  - $R \in \{True, False\}$  as an example

### Probability distributions

- Unobserved random variables have distributions
- A distribution is a TABLE of probabilities of values
- A probability is a single number

### Joint distributions

A *joint distribution* over a set of random variables is a map of assignments to real values

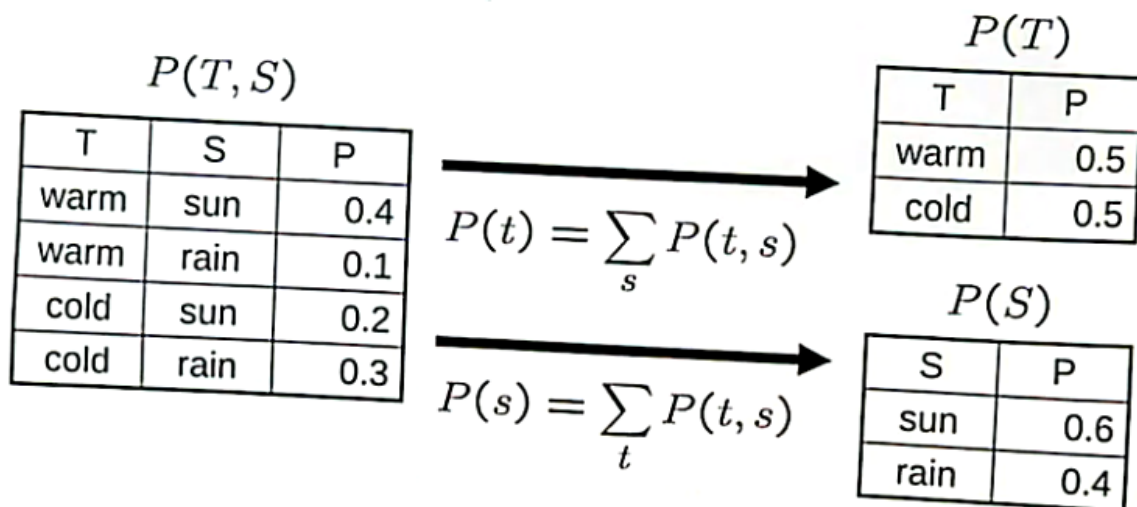
### Events

- A set of assignments
- From a joint distribution we can calculate the probability of any event

### Marginalization

- Marginalization (or summing out) is *projecting* a joint distribution to a sub-distribution over subset of variables.

$$P(X_1 = x_1) = \sum_{x_2} P(X_1 = x_1, X_2 = x_2)$$



### Conditional Probabilities

- A conditional probability is the probability of an event given another event, *center of a venn diagram*

$$P(a|b) = \frac{P(a, b)}{P(b)}$$

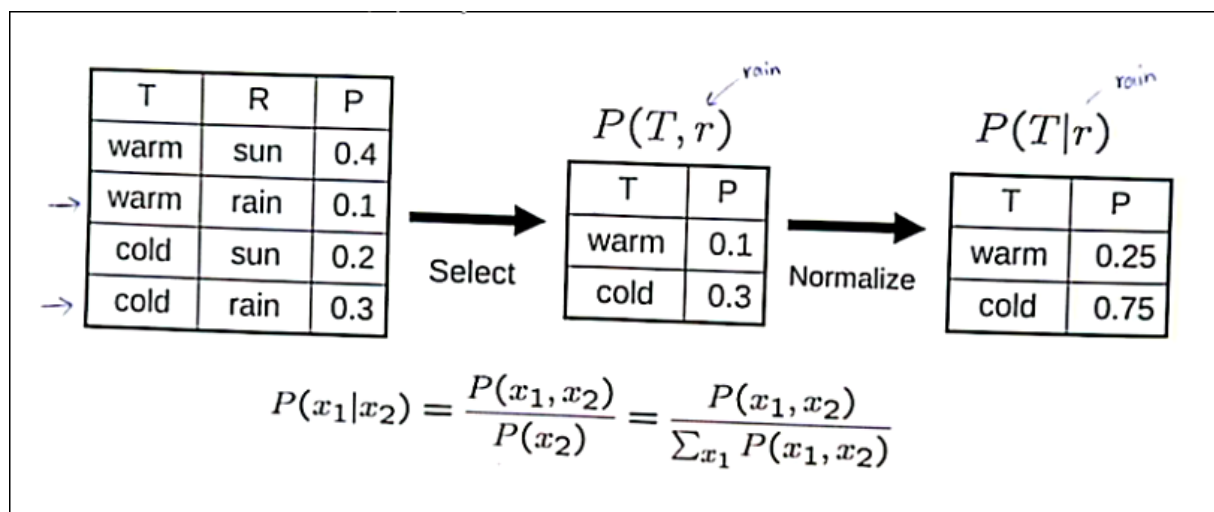
### Conditional Distributions

- Conditional distributions are probability distributions over some variables given fixed values of others.

### Normalization Trick

- A trick to get the whole conditional distribution at once
  - Select the joint probabilities matching the evidence
  - Normalize the selection *divide by total sum so they sum to one*





**Figure 15:** Normalization trick example

### The product rule

- Sometimes joint  $P(X, Y)$  is easy to get
- Sometimes easier to get conditional  $P(X|Y)$

Defined by the following:

$$P(x|y) = \frac{P(x, y)}{P(y)} \equiv P(x, y) = P(x|y) \times P(y)$$

More generally we can write any joint distribution as incremental product of conditional distributions.

$$P(x_1, x_2, x_3) = P(x_1) \times P(x_2|X_1) \times P(x_3|x_1, x_2)$$

$$P(x_1, x_2, \dots, x_n) = \prod_i P(x_i|X_1 \dots x_{i-1})$$

### Probabilistic Inference

- Probabilistic inference: compute a desired probability from other known probabilities *conditional from joint*
- We generally compute conditional probabilities
  - $P(\text{on time} | \text{no report accidents}) = 0.9$
  - These represent the agent's *beliefs* given the evidence

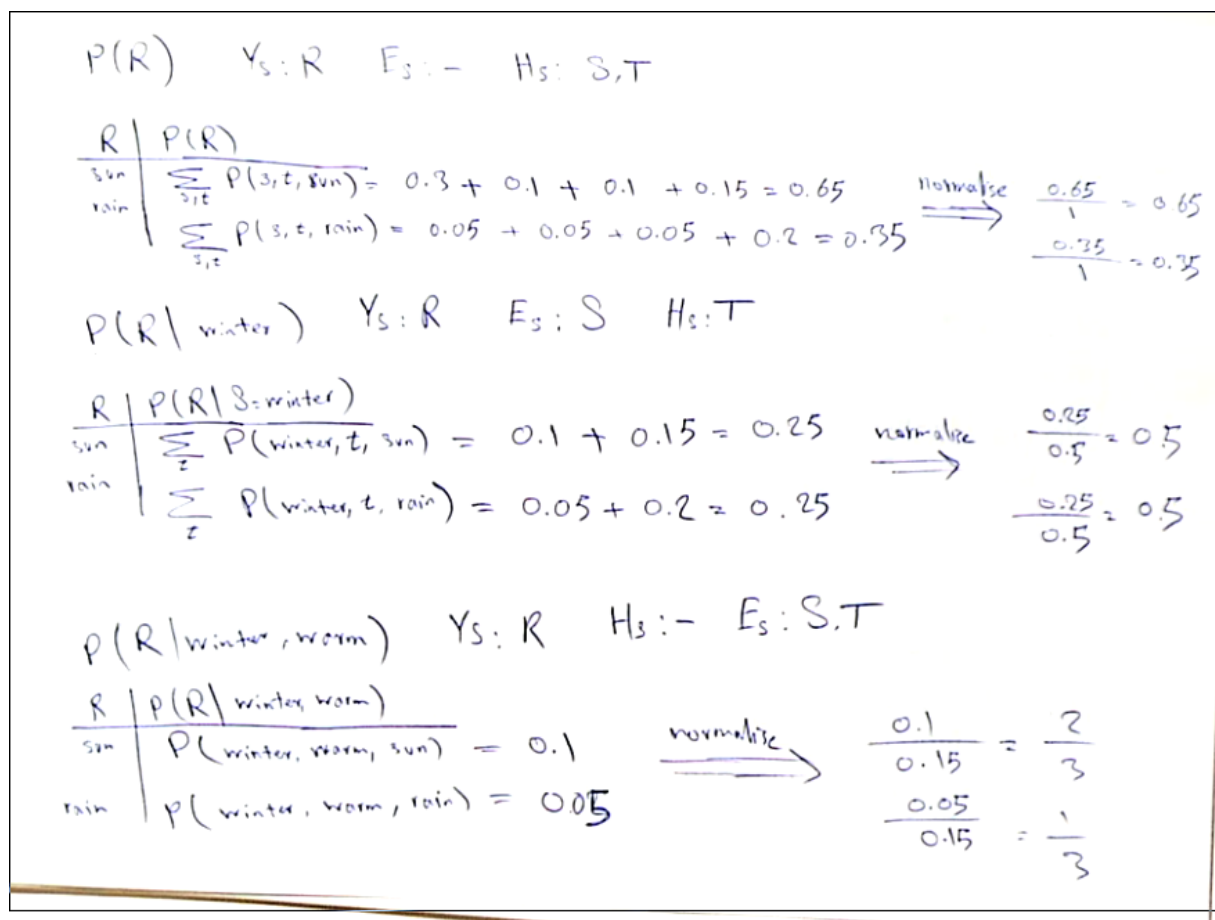
- 
- Probabilities change with new evidence:
    - $P(\text{on time} \mid \text{no accidents, 5 am}) = 0.95$
    - $P(\text{on time} \mid \text{no accidents, 5 am, raining}) = 0.80$
    - Observing new evidence causes beliefs to be updated

### **Inference by Enumeration:**

- we want  $P(Y_1 \dots Y_m \mid e_1 \dots e_k)$ 
  - Evidence variables  $(E_1 \dots E_k) = (e_1 \dots e_k)$
  - Query variables:  $Y_1 \dots Y_m$
  - Hidden variables:  $H_1 \dots H_r$
- First, select the entries consistent with the evidence
- Second, sum out  $H$ :

$$P(Y_1 \dots Y_m, e_1 \dots e_k) = \sum_{h_1 \dots h_r} P(Y_1 \dots Y_m, h_1 \dots h_r, e_1 \dots e_k) = \{X_1, X_2, \dots, X_n\}$$

- Finally normalize the remaining entries
- Obvious problems
  - Worst-case time complexity  $O(d^n)$
  - Space complexity  $O(d^n)$  to store the joint distribution



**Figure 16:** Example

### Complexity of Models

- Engineers and designers are interested in simple and compact models
  - Simple models are easier to build
  - Simple models are easier to explain
  - Compact models take less space
  - Usually implies more efficient computational time

If a probabilistic model has multiple distributions, the number of its free parameters is the sum of the number of free parameters of the tables/distributions.

### Independence

- Two variables are independent if  $P(x, y) = P(x)P(y)$ 
  - This says that their joint distribution factors into a product of two simpler distributions

- 
- We can use independence as a modelling assumption
    - Independence can be simplify assumptions

### **Conditional Independence**

- Absolute/Unconditional independence is very rare
- Conditional independence:
  - $\forall x, y, z : P(x, y|z) = P(x|z)P(y|z)$
  - $\forall x, y, z : P(x|y, z) = P(x|z)$