

ENCE360 Mid-term practice test

Instructions

- **This was the 2018 mid-term test, provided for practice. The total number of marks may be different for 2021.**
- Worth a total of 20 marks
- Contribution to final grade: 20%
- Length: 8 questions
- This is an open book test. Notes, textbooks and online resources may be used.
- This open book test is supervised as a University of Canterbury exam. Therefore, you cannot communicate with anyone other than the supervisors during the test. Anyone using email, Facebook or other forms of communication with others will be removed and score zero for this test.
- Please answer all questions carefully and to the point. Check carefully the number of marks allocated to each question. This suggests the degree of detail required in each answer, and the amount of time you should spend on the question.

The 2021 Mid-term test will be run on the Quiz Server. You will submit all code answers to the quiz server. Each incorrect submission to a question will attract a penalty.

You are strongly advised to submit code even if you haven't got it working and run out of time. We will assess non-working code for potential part-marks.

Question 1: Threads (3 Marks)

In program **one.c**, the child thread, **set_data()** is setting **global_data** to a random number with **rand()**. The child thread, **read_data()** is displaying the value of **global_data**. Use a semaphore to protect the critical region of code accessing this global variable.

You need to initialise the semaphores, **read** and **write** to their respective values, and then use **sem_wait** and **sem_post** to do reading and writing in the critical region.

Your task for this question is to complete the code below (in the empty boxes) using as few lines of code as possible (e.g. no error checking).

You can use the following command to compile **one.c**

```
gcc one.c -o one -lpthread
```

The source (same as below) is in **one.c**: complete it by filling in the blanks.

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <semaphore.h>
#include <unistd.h>
#include <assert.h>

#define NUM_THREADS 10

typedef struct {
    // Read and write semaphores for our channel
```

```

sem_t read;
sem_t write;

// Global shared data
int global_data;

} Channel;

void read_data();
void set_data();
void init_channel();

int main()
{
    pthread_t threads[NUM_THREADS];
    Channel channel;
    int i;

    srand(2018);
    init_channel(&channel);

    for(i = 0; i < NUM_THREADS; i += 2) {
        pthread_create(&threads[i], NULL, (void*)&set_data, &channel);
        pthread_create(&threads[i+1], NULL, (void*)&read_data, &channel);
    }

    // wait for threads to finish before continuing
    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////

    printf("exiting\n");
    exit(0);
}

void set_data(Channel *channel)
{
    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////
    assert(channel->global_data == 0);

    printf("Setting data\t");
    channel->global_data = rand();

    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////
}

void read_data(Channel *channel)
{
    int data;

    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////

    data = channel->global_data;

```

```

channel->global_data = 0;
printf("Data: %d\n", data);

// place your code between the lines of //
////////////////////////////////////////

////////////////////////////////////////

}

void init_channel(Channel *channel)
{

    // Initialise count of the read semaphore to 0 (there's nothing to read yet)
    // place your code between the lines of //
    //////////////////////////////////////////

    //////////////////////////////////////////

    // Note the write semaphore is initialised to 1
    //(channel is empty, free for writing)
    // place your code between the lines of //
    //////////////////////////////////////////

    //////////////////////////////////////////

    channel->global_data = 0;

}

```

Question 2 - Semaphores (1 Mark)

In the context of Question 1, if **global_data** was to become an array that could hold **multiple elements**, and had a known fixed size of 10 elements (i.e. we know the array is **empty** at 0, and the array is **full** at 10 elements), and we can only **read** or **write one** element at a time, do we need **one set** (read & write) of semaphores, or do we need **ten sets** (10x read & write) semaphores to organise concurrency?

Penalties are set at 33.33% per submission.

Select one:

- a) We need ten sets of semaphores. Semaphores act exactly like mutexes and can only hold 0 or 1, so we cannot increment a single semaphore to 10. We just use semaphores to ensure one action (like writing) happens before another action (like reading).
- b) We only need ten single semaphores here. If semaphore #1 is 1, then data can be read from the array at index 0, and if semaphore #1 is 0, then data can be written to the array at position 0.
- c) We only need one set of semaphores. We can sem_post the read semaphore until it hits 10, and sem_wait the write semaphore above 0. Semaphores are not mutexes and can take any number.
- d) We only need one single semaphore for this task. The semaphore can be incremented and decremented more than once. Say, if there are 6 writes to the array, the semaphore is at 6, and can be read from 6 times. The semaphore can be a maximum of 10 and minimum of 0.

Question 3: Pipes (5 Marks)

In program **two.c**, two processes are communicating both ways by using two pipes. One process reads input from the user and handles it. The other process performs some translation of the input and hands it back to the first process for printing.

Your task for this question is to complete the code below (in the empty boxes) using as few lines of code as possible (e.g. no error checking).

You can use the following command to compile **two.c**

```
gcc two.c -o two
```

The source (same as below) is in **two.c**: complete it by filling in the blanks.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <ctype.h>

void translator(int input_pipe[], int output_pipe[])
{
    int c;
    char ch;
    int rc;

    /* first, close unnecessary file descriptors */
    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////

    /* enter a loop of reading from the user_handler's pipe, translating */
    /* the character, and writing back to the user handler.                */
    while (read(input_pipe[0], &ch, 1) > 0) {
        c = ch;
        if (isascii(c) && isupper(c))
            c = tolower(c);
        ch = c;

        /* write translated character back to user_handler. */
        // place your code between the lines of //
        //////////////////////////////////////

        //////////////////////////////////////

        if (rc == -1) {
            perror("translator: write");
            close(input_pipe[0]);
            close(output_pipe[1]);
            exit(1);
        }
    }

    close(input_pipe[0]);
    close(output_pipe[1]);
    exit(0);
}
```

```

void user_handler(int input_pipe[], int output_pipe[])
{
    int c;    /* user input - must be 'int', to recognize EOF (= -1). */
    char ch;  /* the same - as a char. */
    int rc;   /* return values of functions. */

    /* first, close unnecessary file descriptors */
    //////////////////////////////////////
    // place your code between the lines of //

    //////////////////////////////////////

    printf("Enter text to translate:\n");
    /* loop: read input from user, send via one pipe to the translator, */
    /* read via other pipe what the translator returned, and write to */
    /* stdout. exit on EOF from user. */
    while ((c = getchar()) > 0) {
        ch = (char)c;

        /* write to translator */
        //////////////////////////////////////
        // place your code between the lines of //

        //////////////////////////////////////

        if (rc == -1) { /* write failed - notify the user and exit. */
            perror("user_handler: write");
            close(input_pipe[0]);
            close(output_pipe[1]);
            exit(1);
        }

        /* read back from translator */
        //////////////////////////////////////
        // place your code between the lines of //

        //////////////////////////////////////

        c = (int)ch;
        if (rc <= 0) {
            perror("user_handler: read");
            close(input_pipe[0]);
            close(output_pipe[1]);
            exit(1);
        }

        putchar(c);
        if (c=='\n' || c==EOF) break;
    }

    close(input_pipe[0]);
    close(output_pipe[1]);
    exit(0);
}

int main(int argc, char* argv[])
{
    int user_to_translator[2];
    int translator_to_user[2];
    int pid;

```

```

int rc;

rc = pipe(user_to_translator);
if (rc == -1) {
    perror("main: pipe user_to_translator");
    exit(1);
}

rc = pipe(translator_to_user);
if (rc == -1) {
    perror("main: pipe translator_to_user");
    exit(1);
}

pid = fork();

switch (pid) {
case -1:
    perror("main: fork");
    exit(1);
case 0:
    translator(user_to_translator, translator_to_user); /* line 'A' */
    exit(0);
default:
    user_handler(translator_to_user, user_to_translator); /* line 'B' */
}

return 0;
}

```

For example:

Input	Result
I wish YOU a HAPPY new YEAR	Enter text to translate: i wish you a happy new year

Question 4 - Pipes (1 Mark)

If we are using a pipe to communicate between a parent and a child process, with the parent writing data for the child to read, why must we always close the ends of the pipe that we are not using?

Penalty is 33.33% per submission.

Select one:

- The kernel buffers data from the parent to the child, and since the file descriptors are duplicated between the parent and the child, the kernel will not know what process the data is meant for, and neither process will be able to read data.
- The pipe command creates a one way pipe, and data can only flow in one direction. However, C is a flexible programming language, which enables us to decide what way the data will flow. We tell the kernel the direction by closing ends of the pipe.
- Processes read from a pipe until EOF (End of File) is received, and since the file descriptors are duplicated between the parent and the child, if the child does not close the write end, EOF will

never be sent, and the child will read forever.

- d) Pipes are virtual file descriptors which must be opened and closed to use. We open the end we want to use, and close the other to signal the kernel the channel data flows across.

Question 5: Sockets - Server (4 Marks)

Question 5 and Question 7 are about implementing a server and client program which communicate over sockets.

Question 5 relates to completing a socket server program. Only submit the socket *server* in this question.

Your task for this question is to complete the code for **threeserver.c** below (in the empty boxes) using as few lines of code as possible (e.g. no error checking).

You can use the following command to compile **threeServer.c**

```
gcc threeServer.c -o threeServer
```

You can use the following commands to run:

- open two terminals
- in one terminal type **./threeServer 1234**
- in the other terminal type **./fourClient localhost 1234**

Note that to test your answer, you will also need to complete the client from Question 7. Always run **threeServer** first before **fourClient**.

The source (same as below) is in **threeserver.c**: complete it by filling in the blanks.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <netdb.h>
#include <errno.h>
#include <unistd.h>

#define MAXDATASIZE 1024

int main(int argc, char *argv[]) {

    if (argc != 2) {
        fprintf(stderr, "usage: threeServer port-number\n");
        exit(1);
    }

    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////

    struct sockaddr_in sa, caller;
    sa.sin_family = AF_INET;
    sa.sin_addr.s_addr = INADDR_ANY;
    sa.sin_port = htons(atoi(argv[1]));
```

```

// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

socklen_t length = sizeof(caller);
// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

char message[MAXDATASIZE] = "congrats you successfully connected to the
server!";
while (strlen(message) > 0)
{
    int numbytes; // number of bytes of data read from socket

    // send data to the client and then get data back from the client:
    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////

    message[numbytes - 1] = '\0';
}

// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

exit (0);
}

```

Question 6 - Sockets - Server (1 Mark)

What function call lets the server know that a client is connecting? Is it **listen()** or **accept()**?

Penalty is 33.33% per submission.

Select one:

- listen()**. **listen()** tells the kernel that we are interested in connections to this socket, and is the first to know when a connection is incoming, so we can call **accept()** to accept the connection and get a new file descriptor.
- accept()**. **accept()** blocks until a client connects to the server, and then returns a new file descriptor which represents the new connection made. **listen()** simply enables the process to start listening for new connections, but does not actually tell the process that an connection attempt has been made.

- c) `listen()` and `accept()` can both do this. Both `listen()` and `accept()` can determine when a client is connecting, and can let the process know to open a new file descriptor for communication.
- d) Neither. `listen()` is used to tell the kernel to start listening for connections, but doesn't return anything so it cannot tell us if a client is connecting. `accept()` only works if we know there is a connection to be made, as it only processes the last stage of communication by opening a file descriptor.

Question 7: Sockets - Client (4 Marks)

Question 5 and Question 7 are about implementing a server and client program which communicate over sockets.

Question 7 relates to completing a socket client program, **fourclient.c**. Only submit the socket *client* in this question.

Your task for this question is to complete the code below (in the empty boxes) using as few lines of code as possible (e.g. no error checking).

You can use the following command to compile **fourClient.c**

```
gcc fourClient.c -o fourClient
```

You can use the following commands to run:

- open two terminals
- in one terminal type **./threeServer 1234**
- in the other terminal type **./fourClient localhost 1234**

Note to test, you need to also complete the server from Question 5. Always run **threeServer** first before **fourClient**.

The source (same as below) is in **fourclient.c**: complete it by filling in the blanks.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <netdb.h>
#include <unistd.h>

#define MAXDATASIZE 1024

int main(int argc, char *argv[])
{
    if (argc != 3) {
        fprintf(stderr, "usage: fourClient hostname port-number\n");
        exit(1);
    }

    // place your code between the lines of //
    //////////////////////////////////////

    //////////////////////////////////////

    struct addrinfo their_addrinfo; // server address info
    struct addrinfo *their_addr = NULL; // connector's address information
```

```

memset(&their_addrinfo, 0, sizeof(struct addrinfo));
their_addrinfo.ai_family = AF_INET; /* communicate using internet address */
their_addrinfo.ai_socktype = SOCK_STREAM; /* use TCP rather than datagram */
/* get IP address info for the hostname (argv[1]) and port (argv[2]) */
getaddrinfo(argv[1], argv[2], &their_addrinfo, &their_addr);

// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

char buffer[MAXDATASIZE]; //buffer contains data from/to server
int numbytes; // number of bytes of data read from socket
// get data from the server:
// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

while (numbytes > 0)
{
    buffer[numbytes-1] = '\0';
    printf("%s\n", buffer); //print out what you received

    // send data to the server and then get data back from the server:
    // place your code between the lines of //
    ///////////////////////////////////

    ///////////////////////////////////
}
// place your code between the lines of //
////////////////////////////////////

////////////////////////////////////

return 0;
}

```

Question 8 - Sockets - Client (1 Mark)

We see that one of the arguments to **connect()** is the server port number. But hang on, we never set any of the arguments to the clients port number. How does the server know what port to connect back to the client on?

Penalty is 33.33% per submission.

Select one:

- a) The client simply uses the same port number as the server port to receive connections back from the server.
- b) We choose a port number by writing code to `bind()` and `listen()` before any calls to `connect()` are made. We then pass this information to the server with a call to `accept()` to show what port we want to use.

- c) The client opens a communication channel to the server using the server port, and the server simply replies over this channel. The client never has to bind any sockets at all to communicate, so there is no need to worry about this.
- d) The system assigns its own port number itself when `connect()` is called, and simply selects a port which isn't being used.